

Ceng 302
Database Management Systems

Relational Model and Algebra

Prof. Dr. Adnan YAZICI

Department of Computer Engineering,

Middle East Technical University

Ankara

(Fall 2021)

Relational Data Model

- Structure of Relational Data Model
- Constraints of the model (keys)
- Fundamental Relational-Algebra-Operations
- Additional Relational-Algebra-Operations
- Extended Relational-Algebra-Operations
- Null Values

Relational Database: Definitions

- **Relational database:** a set of *relations*
- **Relation:** made up of 2 parts:
 - **Instance** : a *table*, with *rows* and *columns*.
#Rows = *cardinality*, #fields = *degree (arity)*.
 - **Schema** : specifies name of relation, plus name and type of each column.
E.g. **Students** (**sid**: string, **name**: string, **login**: string, **age**: integer, **gpa**: real).
- We can think of a relation as a *set* of *rows* or *tuples* (i.e., all rows are distinct).

Basic Structure

- Formally, given sets D_1, D_2, \dots, D_n a **relation** r is a subset of

$$D_1 \times D_2 \times \dots \times D_n$$

Thus, a relation is a set of n -tuples (a_1, a_2, \dots, a_n) where each $a_i \in D_i$

- Example: If
 - $customer_name = \{\text{Jones, Smith, Curry, Lindsay, ...}\}$ /*set of all customers*/
 - $customer_street = \{\text{Main, North, Park, ...}\}$ /* set of all street names */
 - $customer_city = \{\text{Harrison, Rye, Pittsfield, ...}\}$ /*set of all city names*/

Then $r = \{(\text{Jones, Main, Harrison}), (\text{Smith, North, Rye}), (\text{Curry, North, Rye}), (\text{Lindsay, Park, Pittsfield})\}$ is a relation over

customer_name x customer_street x customer_city

Attribute Types

- Each attribute of a relation has a name
- The set of allowed values for each attribute is called the **domain** of the attribute
- Attribute values are (normally) required to be **atomic**; that is, indivisible
 - E.g. the value of an attribute can be an account number, but cannot be a set of account numbers
- Domain is said to be atomic if all its members are atomic
- The special value *null* is a member of every domain
- The null value causes complications in the definition of many operations

Relation Schema

- A_1, A_2, \dots, A_n are *attributes*
- $R = (A_1, A_2, \dots, A_n)$ is a *relation schema*

Example:

$Customer_schema = (customer_name, customer_street, customer_city)$

- $r(R)$ denotes a *relation* r on the *relation schema* R

Example:

$customer (Customer_schema)$

Relation Instance

- The current values (*relation instance*) of a relation are specified by a table
- An element t of r is a *tuple*, represented by a *row* in a table

Customer

<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>
<i>Jones</i>	<i>Main</i>	<i>Harrison</i>
<i>Smith</i>	<i>North</i>	<i>Rye</i>
<i>Curry</i>	<i>North</i>	<i>Rye</i>
<i>Lindsay</i>	<i>Park</i>	<i>Pittsfield</i>

Example Instance of Students Relation

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

- Cardinality = 3, degree = 5, all rows distinct
- Do all columns in a relation instance have to be distinct?

Database

- A database consists of multiple relations
- Information about an enterprise is broken up into parts, with each relation storing one part of the information

account : stores information about accounts

depositor : stores information about which customer owns which account

customer : stores information about customers

- Storing all information as a single relation such as
bank (account_number, balance, customer_name, ..) results in
 - repetition of information
 - the need for null values
- **Normalization theory** deals with how to design relational schemas

Relational Query Languages

- A major strength of the relational model:
 - supports simple, powerful *querying* of data.
- Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
 - Precise semantics for relational queries.
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.

The SQL Query Language

- Developed by IBM (system R) in the 1970s
- Need for a standard since it is used by many vendors
- Standards:
 - SQL-86
 - SQL-89 (minor revision)
 - SQL-92 (major revision, current standard)
 - SQL-99 (major extensions)

Creating Relations in SQL

- Creates the Students relation. Observe that the type (**domain**) of each field is specified and enforced by the DBMS whenever tuples are added or modified.

```
CREATE TABLE Students  
(sid: CHAR(20),  
 name: CHAR(20),  
 login: CHAR(10),  
 age: INTEGER,  
 gpa: REAL)
```

- As another example, the Enrolled table holds information about courses that students take.

```
CREATE TABLE Enrolled  
(sid: CHAR(20),  
 cid: CHAR(20),  
 grade: CHAR(2))
```

Adding and Deleting Tuples

- Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)  
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```

- Can delete all tuples satisfying some condition (e.g., *name* = Smith):

```
DELETE  
FROM Students S  
WHERE S.name = 'Smith'
```

* *Powerful variants of these commands are available.*

Modifying Tuples

- Can modify the column values in an existing row using:

```
UPDATE Students S
SET      S.age = S.age + 1, S.gpa = S.gpa - 1
WHERE    S.sid = 53688
```

```
UPDATE Students S
SET      S.gpa = S.gpa - 0.1
WHERE    S.gpa >= 3.3
```

Integrity Constraints (ICs)

- **IC:** condition that must be true for *any* instance of the database; e.g., *domain constraints*.
 - ICs are specified when schema is defined.
 - ICs are checked when relations are modified.
- A *legal* instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances.
- If the DBMS checks ICs, stored data is more faithful to real-world meaning.
 - IC Integrity constraints ensure that the data insertion, updating, and other processes have to be performed in such a way that data integrity is not affected.
 - ICs are used to guard against accidental damage to the database.
 - ICs mainly avoid data entry errors.

Keys

- Let $K \subseteq R$
- K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each relation $r(R)$.
 - **Example:** $\{customer_name, customer_street\}$ and $\{customer_name\}$
are both superkeys of the *Customer* relation, if no two customers can possibly have the same name.
 - In real life, an attribute such as *customer_id* would be used instead of *customer_name* to uniquely identify customers.

Keys (Cont.)

- K is a **candidate key** if K is minimal
Example: $\{customer_name\}$ is a candidate key for *Customer*, since it is a superkey and no subset of it is a superkey.
- **Primary key:** a candidate key is chosen as the principal means of identifying tuples within a relation
 - Should we choose an attribute whose value never, or very rarely, changes.
 - E.g. *email address* is unique, but it may change

Examples:

- *sid* is a *key* for Students. (What about *name*?)
- *email* is also a candidate key for students.
- The set $\{sid, gpa\}$ is a *superkey*.

Primary and Candidate Keys in SQL

- Possibly many candidate keys (specified using UNIQUE), one of which is chosen as the *primary key*.
- “For a given student and course, there is a single grade.”

sid	cid	grade
53666	COP4	A
53666	COP4	B-
53666	CDA3	A

```
CREATE TABLE Enrolled
(sid CHAR(20)
 cid CHAR(20),
 grade CHAR(2),
PRIMARY KEY (sid,cid))
```

Primary and Candidate Keys in SQL

- Possibly many *candidate keys* (specified using **UNIQUE**), one of which is chosen as the *primary key*.
- “Students can take only one course and receive a single grade for that course; further, no two students in a course receive the same grade.”

sid	cid	grade
53666	COP4	A
53666	CDA3	B-
53444	COP4	A

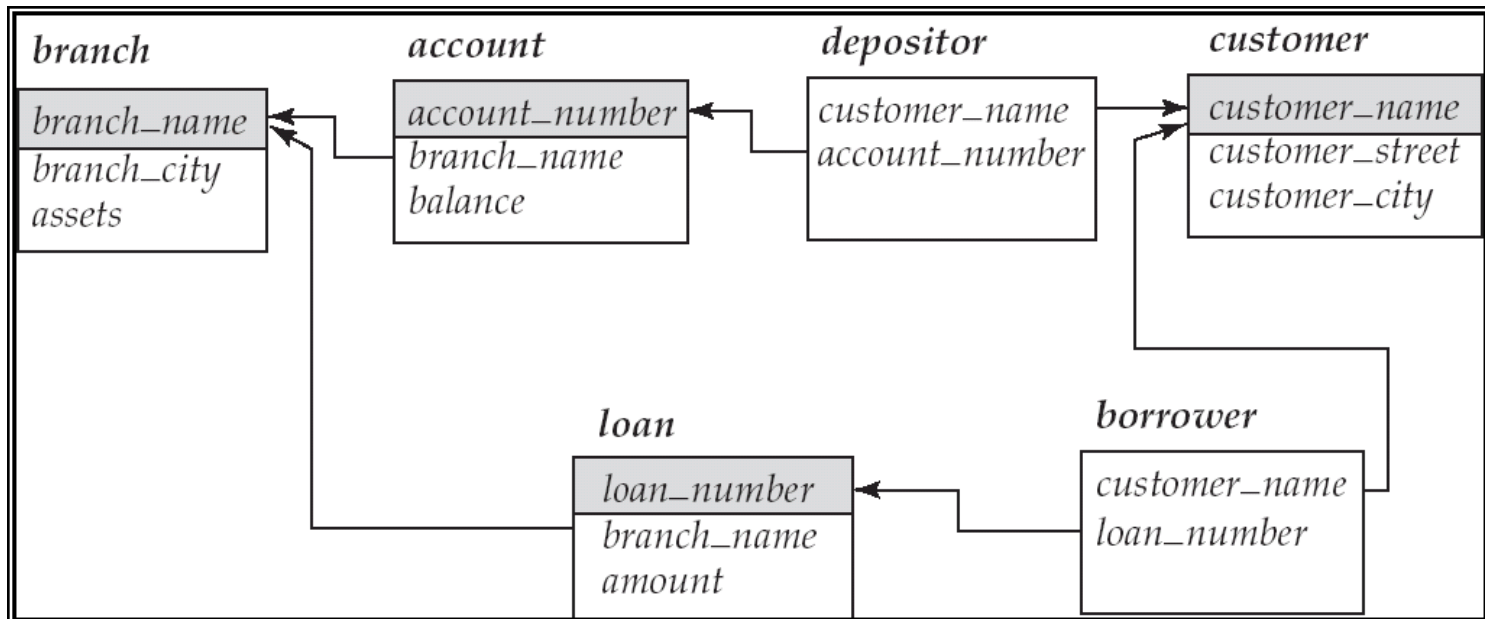
```
CREATE TABLE Enrolled
(sid CHAR(20)
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid),
UNIQUE (cid, grade) )
```

Foreign Keys, Referential Integrity

- *Foreign Key (FK)*: Set of fields in one relation that is used to 'refer' to a tuple in another relation. (FK must correspond to primary key of the second relation.) Like a 'logical pointer'.
- E.g. *sid* in the *Enrolled* relation is a foreign key referring to *Students*:
 - Enrolled(*sid*: string, *cid*: string, *grade*: string)
 - If all foreign key constraints are enforced, *referential integrity* is achieved, i.e., no dangling references.

Foreign Keys

- A relation schema may have an attribute (X) that corresponds to the **primary key** of another relation. This attribute X is called a **foreign key**.
 - E.g. *customer_name* and *account_number* attributes of *depositor* are foreign keys to *customer* and *account* respectively.
 - Only values occurring in the primary key attribute of the **referenced relation** (*customer*) may occur as the foreign key attribute of the **referencing relation** (*depositor*).
- Schema diagram



Foreign Keys in SQL

- Only students listed in the **Students** relation should be allowed to enroll for courses. This is called the **referential integrity constraint**.


```
CREATE TABLE Enrolled  
  (sid CHAR(20), cid CHAR(20), grade CHAR(2),  
   PRIMARY KEY (sid,cid),  
   FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8



Enforcing Referential Integrity

- Consider *Students* and *Enrolled*; *sid* in *Enrolled* is a foreign key that references *Students*.
- What should be done if an *Enrolled* tuple with a non-existent *student id* is inserted? (*Reject it!*)
- What should be done if a *Students* tuple is deleted?
 - Also delete all *Enrolled* tuples that refer to it.
 - Disallow deletion of a *Students* tuple that is referred to.
 - Set *sid* in *Enrolled* tuples that refer to it to a *default sid*.
 - (In SQL, also: Set *sid* in *Enrolled* tuples that refer to it to a special value *null*, denoting 'unknown' or 'inapplicable'.)
- Similar if primary key of *Students* tuple is updated.

Referential Integrity in SQL/92

➤ SQL/92 supports all 4 options on deletes and updates.

- Default is **NO ACTION**
(*delete/update is rejected*)
- **CASCADE** (also delete all tuples that refer to deleted tuple)
- **SET NULL / SET DEFAULT**
(sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20) DEFAULT '9999',
 cid CHAR(20),
 grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE NO ACTION )
```


Query Languages

- Language in which user requests information from the database.
- Categories of languages
 - Procedural
 - Non-procedural, or declarative
- “Pure” languages:
 - Relational algebra
 - Tuple relational calculus
 - Domain relational calculus
- Pure languages form underlying basis of query languages that people use.

Relational Algebra

- The Relational Algebra is procedural; you tell it how to construct the result
- It consists of a set of **operators** which, when applied to relations, yield relations (closed algebra)

$R \cup S$	union , R UNION S
$R \cap S$	intersection , R INTERSECT S
$R \setminus S$	set difference , R MINUS S
$R \times S$	Cartesian product , R JOIN S (no shared attributes)
$\pi_{A_1, A_2, \dots, A_n}(R)$	projection , R[A1, A2, ..., An]
$\sigma_{\text{expression}}(R)$	selection , R WHERE EXPRESSION
$R \bowtie S$	natural join , R JOIN S (no shared attributes)
$R \bowtie_{\theta} S$	theta-join , via selection from \times
$R \div S$	divideby , R DIVIDEBY S
$\rho [A_1 B_1, \dots, A_n B_n]$	rename , R[A1 B1, ..., An Bn]

Relational Algebra

- **Basic operations:**

- Selection (σ) Selects a subset of rows from relation.
- Projection (π) Deletes unwanted columns from relation.
- Cross-product (\times) Allows us to combine two relations.
- Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
- Union (\cup) Tuples in reln. 1 and in reln. 2.

or

- Intersect (\cap) Tuples in reln. 1 and in reln. 2.

- **Additional operations:**

- Intersection (or union), join, division, renaming: Not essential, but (very!) useful.
- Since each operation returns a relation, operations can be *composed*! (Algebra is “closed”.)

Select Operation – Example

Relation r:

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

$\sigma_{A=B \wedge D > 5}(r)$:

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
α	α	1	7
β	β	23	10

Select Operation

- Notation: $\sigma_p(r)$
- p is called the **selection predicate**
- Defined as:

$$\sigma_p(r) = \{t \mid t \in r \text{ and } p(t)\}$$

Where p is a formula in propositional calculus consisting of **terms** connected by \wedge (**and**), \vee (**or**), \neg (**not**)

Each **term** is one of:

$\langle \text{attribute} \rangle \text{ op } \langle \text{attribute} \rangle$ or $\langle \text{constant} \rangle$

where **op** is one of: $=, \neq, >, \geq, <, \leq$

- Example of selection:

$$\sigma_{\text{branch_name} = \text{"Perryridge"}}(\text{account})$$

Project Operation – Example

Relation r :

A	B	C
α	10	1
α	20	1
β	30	1
β	40	2

$\Pi_{A,C}(r)$:

A	C
α	1
α	1
β	1
β	2

 $=$

A	C
α	1
β	1
β	2

Project Operation

- Notation:

$$\Pi_{A_1, A_2, \dots, A_k}(r)$$

where A_1, A_2 are attribute names and r is a relation name.

- The result is defined as the relation of k columns obtained by erasing the columns that are not listed
- Duplicate rows removed from result, since relations are sets
- **Example:** To eliminate the *branch_name* attribute of *account*

$$\Pi_{\text{account_number, balance}}(\text{account})$$

Union Operation – Example

Relations r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

$r \cup s$:

A	B
α	1
α	2
β	1
β	3

Union Operation

- Notation: $r \cup s$
- Defined as:

$$r \cup s = \{t \mid t \in r \text{ or } t \in s\}$$

- For $r \cup s$ to be valid, they must be **union compatible**, which means;
 1. r, s must have the *same* **arity** (same number of attributes)
 2. The attribute domains must be **compatible**

Example: 2nd column of r deals with the same type of values as does the 2nd column of s)

Example: to find all customers with either an account or a loan

$$\Pi_{customer_name}(depositor) \cup \Pi_{customer_name}(borrower)$$

Set-Intersection Operation – Example

Relation r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

$r \cap s$

A	B
α	2

Notation: $r \cap s$

Defined as:

$$r \cap s = \{ t \mid t \in r \textbf{ and } t \in s \}$$

Assume:

r, s are **union compatible**

Note: $r \cap s = r - (r - s)$

Set Difference Operation – Example

Relations r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

$r - s$:

A	B
α	1
β	1

Set Difference Operation

- Notation $r - s$
- Defined as:

$$r - s = \{t \mid t \in r \text{ and } t \notin s\}$$

- Set differences must also be **union compatible** relations.
- As might noticed, **Difference**, **Union** and **Intersect** operations must be ***union-compatible***;
 - Same number of fields.
 - Attribute domains of r and s must be compatible, that is 'Corresponding' fields have the same type.

Cartesian-Product Operation – Example

Relations r, s :

A	B
-----	-----

α	1
β	2

r

C	D	E
-----	-----	-----

α	10	a
β	10	a
β	20	b
γ	10	b

s

$r \times s$:

A	B	C	D	E
-----	-----	-----	-----	-----

α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

Cartesian-Product Operation

- Notation $r \times s$
- Defined as:

$$r \times s = \{t \ q \mid t \in r \textbf{ and } q \in s\}$$

- Assume that attributes of $r(R)$ and $s(S)$ are disjoint. That is,

$$R \cap S = \emptyset.$$

- If attributes of $r(R)$ and $s(S)$ are not disjoint, then renaming must be used.

Composition of Operations

- Can build expressions using multiple operations
- Example: $\sigma_{A=C}(r \times s)$
- $r \times s$

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
α	1	α	10	<i>a</i>
α	1	β	10	<i>a</i>
α	1	β	20	<i>b</i>
α	1	γ	10	<i>b</i>
β	2	α	10	<i>a</i>
β	2	β	10	<i>a</i>
β	2	β	20	<i>b</i>
β	2	γ	10	<i>b</i>

- $\sigma_{A=C}(r \times s)$

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>
α	1	α	10	<i>a</i>
β	2	β	10	<i>a</i>
β	2	β	20	<i>b</i>

Natural-Join Operation

Notation: $r \bowtie s$

- Let r and s be relations on schemas R and S respectively.
Then, $r \bowtie s$ is a relation on schema $R \cup S$ obtained as follows:
 - Consider each pair of tuples t_r from r and t_s from s .
 - If t_r and t_s have the same value on each of the attributes in $R \cap S$, add a tuple t to the result, where
 - t has the same value as t_r on r
 - t has the same value as t_s on s
- Example:
 - $R = (A, B, C, D)$
 - $S = (E, B, D)$
 - Result schema = (A, B, C, D, E)
 - $r \bowtie s$ is defined as:

$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$$

Natural Join Operation – Example

- Relations r , s :

A	B	C	D
α	1	α	a
β	2	γ	a
γ	4	β	b
α	1	γ	a
δ	2	β	b

r

B	D	E
1	a	α
3	a	β
1	a	γ
2	b	δ
3	b	ϵ

s

$r \bowtie s$

A	B	C	D	E
α	1	α	a	α
α	1	α	a	γ
α	1	γ	a	α
α	1	γ	a	γ
δ	2	β	b	δ

$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B=s.B \wedge r.D=s.D} (r \times s))$$

Rename Operation

- Allows us to name, and therefore to refer to, the results of relational-algebra expressions.
- Allows us to refer to a relation by more than one name.
- Example:

$$\rho_X(E)$$

returns the expression E under the name X

- If a relational-algebra expression E has arity n , then

$$\rho_{X(A_1, A_2, \dots, A_n)}(E)$$

returns the result of expression E under the name X , and with the attributes renamed to A_1, A_2, \dots, A_n .

Division Operation

- Notation: $r \div s$
- Suited to queries that include the phrase “**for all**”.
- Let r and s be relations on schemas R and S respectively where
 - $R = (A_1, \dots, A_m, B_1, \dots, B_n)$
 - $S = (B_1, \dots, B_n)$

The result of $r \div s$ is a relation on schema

$$R - S = (A_1, \dots, A_m)$$

$$r \div s = \{ t \mid t \in \Pi_{R-S}(r) \wedge \forall u \in s (tu \in r) \}$$

where tu means the concatenation of tuples t and u to produce a single tuple.

- $r \div s$ contains all t tuples such that for every u tuple in s , there is a tu tuple in r . Or if the set of u values associated with a t value in r contains all u values in s , the t value is in $r \div s$.

Division Operation – Example

Relations r, s :

A	B
α	1
α	2
α	3
β	1
γ	1
δ	1
δ	3
δ	4
ϵ	6
ϵ	1
β	2

r

B
1
2

s

$r \div s$:

A
α
β

$$R - S = (A_1, \dots, A_m)$$

$$r \div s = \{ t \mid t \in \Pi_{R-S}(r) \wedge \forall u \in s (tu \in r) \}$$

Examples of Division R/S_i

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

R

pno
p2

S_1

sno
s1
s2
s3
s4

R/S_1

pno
p2
p4

S_2

sno
s1
s4

R/S_2

pno
p1
p2
p4

S_3

sno
s1

R/S_3

Another Division Example

Relations r, s :

A	B	C	D	E
α	a	α	a	1
α	a	γ	a	1
α	a	γ	b	1
β	a	γ	a	1
β	a	γ	b	3
γ	a	γ	a	1
γ	a	γ	b	1
γ	a	β	b	1

r

D	E
a	1
b	1

s

$r \div s$:

A	B	C
α	a	γ
γ	a	γ

$$R - S = (A_1, \dots, A_m)$$

$$r \div s = \{ t \mid t \in \Pi_{R-S}(r)$$

$$\wedge \forall u \in s (tu \in r) \}$$

Division Operation (Cont.)

Property

- Let $q = r \div s$
- Then q is the largest relation satisfying $q \times s \subseteq r$

- Definition in terms of the basic algebra operation

Let $r(R)$ and $s(S)$ be relations, and let $S \subseteq R$


$$r \div s = \Pi_{R-S}(r) - \Pi_{R-S}((\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r))$$

- $\Pi_{R-S,S}(r)$ simply reorders attributes of r
- $\Pi_{R-S}(\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r)$ gives those tuples t in $\Pi_{R-S}(r)$ such that for some tuple $u \in s$, $tu \notin r$.

Division Operation (Cont.)

$$r \div s = \Pi_{R-S}(r) - \Pi_{R-S} [(\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r)]$$

- $\Pi_{R-S,S}(r)$ simply reorders attributes of r
- $\Pi_{R-S}(\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r)$ gives those tuples t in $\Pi_{R-S}(r)$ such that for some tuple $u \in s$, $tu \notin r$.



Example Instances

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
22	103	12/10/96
22	102	13/10/96

Sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

- We will use these instances of the **Sailors**, **Boats** and **Reserves** relations in our examples.

Boats

<u>bid</u>	bname	color
101	Intertake	blue
102	Intertake	red
103	Clipper	green

Query: "Find sailors who've reserved all boats."

Division in SQL

Query: Find sailors who've reserved **all** boats.

- Let's do it without EXCEPT:

```
SELECT S.sname
FROM Sailors S
```

Sailors S such that ...

```
WHERE NOT EXISTS (SELECT B.bid
                   FROM Boats B
```

there is no boat B without

```
WHERE NOT EXISTS (SELECT R.bid
                   FROM Reserves R
```

a Reserves tuple showing S reserved B

```
WHERE R.bid=B.bid
      AND R.sid=S.sid))
```

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
      ((SELECT B.bid
        FROM Boats B)
EXCEPT
 (SELECT R.bid
  FROM Reserves R
 WHERE R.sid=S.sid))
```

or “Select each sailor such that there does not exist a boat that the sailor does not reserve it.”

Querying using Relational Algebra

Query: Find names of sailors who've reserved a red boat

- Information about boat color only available in Boats; so need an extra join:

$$\pi_{sname}((\sigma_{color='red'} Boats) \bowtie Reserves \bowtie Sailors)$$

- Alternative solution:

$$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color='red'} Boats) \bowtie Res) \bowtie Sailors)$$

A query optimizer can find this, given the first solution!

Assignment Operation

- The assignment operation (\leftarrow) provides a convenient way to express complex queries.
 - Write query as a sequential program consisting of
 - a series of assignments
 - followed by an expression whose value is displayed as a result of the query.
 - Assignment must always be made to a temporary relation variable.

Example: Write $r \div s$ as

$$temp1 \leftarrow \Pi_{R-S}(r)$$

$$temp2 \leftarrow \Pi_{R-S}((temp1 \times s) - \Pi_{R-S,S}(r))$$

$$result = temp1 - temp2$$

- The result to the right of the \leftarrow is assigned to the relation variable on the left of the \leftarrow .
- May use variable in subsequent expressions.

Querying using Relational Algebra

Query: Find names of sailors who've reserved boat #103

❖ Solution 1: $\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$

❖ Solution 2: $\rho(Temp1, \sigma_{bid=103} Reserves)$

$\rho(Temp2, Temp1 \bowtie Sailors)$

$\pi_{sname}(Temp2)$

❖ Solution 3: $\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$

Aggregate Functions and Operations

- **Aggregation function** takes a collection of values and returns a single value as a result.

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

- **Aggregate operation** in relational algebra

$$G_1, G_2, \dots, G_n \mathcal{G}_{F_1(A_1), F_2(A_2), \dots, F_n(A_n)}(E)$$

E is any relational-algebra expression

- G_1, G_2, \dots, G_n is a list of attributes on which to group (can be empty)
- Each F_i is an aggregate function
- Each A_i is an attribute name

Aggregate Operation – Example

Relation r :

A	B	C
α	α	7
α	β	7
β	β	3
β	β	10

$g_{\text{sum}(C)}(r)$

$\text{sum}(C)$
27

Aggregate Operation – Example

Relation ***account*** grouped by ***branch-name***:

<i>branch_name</i>	<i>account_number</i>	<i>balance</i>
Perryridge	A-102	400
Perryridge	A-201	900
Brighton	A-217	750
Brighton	A-215	750
Redwood	A-222	700

branch_name \mathcal{G} **sum**(*balance*) (*account*)

<i>branch_name</i>	sum (<i>balance</i>)
Perryridge	1300
Brighton	1500
Redwood	700

Aggregate Functions (Cont.)

Result of aggregation does not have a name

- Can use rename operation to give it a name
- For convenience, we permit renaming as part of aggregate operation

branch_name \mathcal{G} *sum(balance) as sum_balance* (*account*)

Outer Join

- An extension of the **join** operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values:
 - *null* signifies that the value is **unknown** or **does not exist**
 - All comparisons involving *null* are (roughly speaking) **false** by definition.
 - We shall study precise meaning of comparisons with nulls later

Outer Join – Example

Relation *loan*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

Relation *borrower*

<i>customer_name</i>	<i>loan_number</i>
Jones	L-170
Smith	L-230
Hayes	L-155

Outer Join – Example

loan

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

borrower

<i>customer_name</i>	<i>loan_number</i>
Jones	L-170
Smith	L-230
Hayes	L-155

Join

loan ⋈ *borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

Left Outer Join

loan ⋈_L *borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>

Outer Join – Example

loan

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

borrower

<i>customer_name</i>	<i>loan_number</i>
Jones	L-170
Smith	L-230
Hayes	L-155

Right Outer Join

loan ⋈_r *borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	<i>null</i>	<i>null</i>	Hayes

Full Outer Join

loan ⋈_f *borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>
L-155	<i>null</i>	<i>null</i>	Hayes

Relational Model - Operations

Powerful set-oriented query languages:

- **Relational Algebra**: **procedural**; describes how to compute a query

Operators; Union, Select, Project, Cartesian Product, Difference, and

Intersect, Join, Division, Outer Join, Outer Union, etc.

- **Relational Calculus**: **declarative**; describes the desired result, Insert, delete, and update capabilities

e.g., SQL, QBE

Outer Union

Display all data values from Table df1 and table df4, but overlay common attributes.

df1					Result					
	A	B	C	D		A	B	C	D	F
0	A0	B0	C0	D0	0	A0	B0	C0	D0	NaN
1	A1	B1	C1	D1	1	A1	B1	C1	D1	NaN
2	A2	B2	C2	D2	2	A2	B2	C2	D2	NaN
3	A3	B3	C3	D3	3	A3	B3	C3	D3	NaN
df4					2	NaN	B2	NaN	D2	F2
	B	D	F		3	NaN	B3	NaN	D3	F3
2	B2	D2	F2	6	NaN	B6	NaN	D6	F6	
3	B3	D3	F3	7	NaN	B7	NaN	D7	F7	
6	B6	D6	F6							
7	B7	D7	F7							

Outer Union

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *  
  from one  
outer union corr  
select *  
  from two;
```

Final Results

X	A	B
1	a	
1	a	
1	b	
2	c	
3	v	
4	e	
6	g	
1		x
2		y
3		z
3		v
5		w

Relational Calculus

- The **Relational Calculus** is non-procedural. It allows you to express a result relation using a predicate on tuple variables (**tuple calculus**):

$$\{ t \mid P(t) \}$$

or on domain variables (**domain calculus**):

$$\{ \langle x_1, x_2, \dots, x_n \rangle \mid P(\langle x_1, x_2, \dots, x_n \rangle) \}$$

- You tell the system which result you want, but not how to construct it.

Relational Calculus

FLT-WEEKDAY

<u>flt#</u>	<u>weekday</u>
-------------	----------------

Query: Find FLT# for all flights scheduled for Mondays

Tuple calculus:

$\{t.FLT\# \mid FLT-WEEKDAY(t) \wedge t.WEEKDAY = MO\}$

Domain calculus:

$\{ \langle FLT\# \rangle \mid \langle FLT\#, WEEKDAY \rangle \in FLT-WEEKDAY \wedge WEEKDAY = MO \}$

Relational Calculus

Tuple and **domain calculus** for **join** operation: $\{t \mid P(t)\}$

FLT-WEEKDAY

<u>flt#</u>	weekday
-------------	---------

FLT-INSTANCE

<u>flt#</u>	<u>date</u>	plane#	#avail-seats
-------------	-------------	--------	--------------

Query: Find and make a list with complete flight instance information

Tuple Calculus:

$\{s.FLT\#, s.WEEKDAY, t.DATE, t.PLANE\#, t.\#AVAIL-SEATS \mid FLT-WEEKDAY(s) \wedge FLT-INSTANCE(t) \wedge s.FLT\# = t.FLT\# \}$

Domain Calculus:

$\{ \langle FLT\#, WEEKDAY, DATE, PLANE\#, \#AVAIL-SEATS \rangle \mid \langle FLT\#, WEEKDAY \rangle \in FLT-WEEKDAY \wedge \langle FLT\#, DATE, PLANE\#, \#AVAIL-SEATS \rangle \in FLT-INSTANCE \wedge FLT\# = FLT\# \}$

Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an **unknown** value or that a value **does not exist**.
- The result of any arithmetic expression involving null is *null*.
- **Aggregate functions** simply ignore **null** values (as in SQL)
- For duplicate elimination and grouping, null is treated like any other value, and two nulls are assumed to be the same (as in SQL)

Null Values

- Comparisons with null values return the special truth value: ***unknown***
 - If *false* were used instead of *unknown*, then *not* ($A < 5$) would not be equivalent to $A \geq 5$
- Three-valued logic using the truth value ***unknown***:
 - OR: (*unknown* **or** *true*) = *true*,
(*unknown* **or** *false*) = *unknown*
(*unknown* **or** *unknown*) = *unknown*
 - AND: (*true* **and** *unknown*) = *unknown*,
(*false* **and** *unknown*) = *false*,
(*unknown* **and** *unknown*) = *unknown*
 - NOT: (**not** *unknown*) = *unknown*
- In SQL “*P* is ***unknown***” evaluates to true if predicate *P* evaluates to *unknown*
- Result of select predicate is treated as *false* if it evaluates to *unknown*