

**Ceng 302**  
**Database Management Systems**

**SQL: Structured Query Language**

**Prof. Dr. Adnan YAZICI**  
Department of Computer Engineering,  
Middle East Technical University  
(Fall 2021)

# DML - Overview

- select-from-where
- set operations
- nested subqueries
- built-in functions
- insert, delete, update
- joins and outer joins
- group by and having clauses
- ordering of rows
- embedded DML
- views
- indexing
- triggers
- integrity constraints
- db authorizations
- recursive queries

# Relational Query Languages

- *Query languages*: Allow manipulation and retrieval of data from a database.
- Relational model supports simple, powerful QLs:
  - Strong formal foundation based on formal logic.
  - Allows for much optimization.
- Query Languages **!=** programming languages!
  - QLs are not expected to be “Turing complete” or “computationally universal.”
  - QLs are not intended to be used for complex calculations.
  - QLs support easy, efficient access to large data sets.

## Interactive DML - select-from-where

```
SELECT A1, A2, ... An  
FROM   R1 , R2 , ... Rm  
WHERE  P
```

- the **SELECT** clause specifies the columns of the result
- the **FROM** clause specifies the tables to be scanned in the query
- the **WHERE** clause specifies the condition on the columns of the tables in the **FROM** clause
- equivalent algebra statement:

$$\pi_{A_1, A_2, \dots, A_n} ( \sigma_P ( R_1 \times R_2 \times \dots \times R_m ) )$$

# Basic SQL Query

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

- *relation-list* A list of relation names (possibly with a *range-variable* after each relation name).
- *target-list* A list of attributes of relations included in *relation-list*
- *qualification* Comparisons (**Attr** *op* **const** or **Attr1** *op* **Attr2**, where *op* is one of ( <, >, =, ≤, ≥, ≠ ) combined using **AND**, **OR** and **NOT**.
- **DISTINCT** is an optional keyword indicating that the answer should not contain duplicates. Default is that duplicates are **not** eliminated!

# Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
  - Compute the cross-product of *relation-list*.
  - Discard resulting tuples if they fail *qualifications*.
  - Delete attributes that are not in *target-list*.
  - If **DISTINCT** is specified, eliminate duplicate rows.
- This strategy is probably the least efficient way to compute a query! An optimizer will find more efficient strategies to compute the same answers.

## Example Instances

*Reserves*

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

*Sailors*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

- We use these instances of the **Sailors**, **Boats** and **Reserves** relations in our examples.

*Boats*

<u>bid</u>	bname	color
101	Intertake	blue
102	Intertake	red
103	Clipper	green
104	Marine	red

# Example of Conceptual Evaluation

**Query:** Retrieve the sailors names who reserved the boat id 103.

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND R.bid=103
```

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96



# A Note on Range Variables

- Needed only if the same relation appears twice in the FROM clause. The previous query can also be written as:

```
SELECT sname
FROM   Sailors, Reserves
WHERE  Sailors.sid=Reserves.sid
      AND bid=103
```

OR

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND bid=103
```

*Using range variables is a good style!*

# Expressions and Strings

*Query: 'Find triples (of ages of sailors and two fields defined by expressions) for sailors whose names begin and end with B and contain at least three characters.'*

```
SELECT S.age, age1 = S.age-5, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname LIKE 'B_%B'
```

This query illustrates the use of arithmetic expressions and string pattern matching:

- **AS** and **=** are two ways to name fields in result.
- **LIKE** is used for string matching. **`\_`** stands for any one character and **`%`** stands for 0 or more arbitrary characters.

# Nested Queries

*Q: “Find names of sailors who’ve reserved boat #103”*

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid IN (SELECT R.sid
                 FROM   Reserves R
                 WHERE  R.bid=103)
```

- **IN** operator performs a direct match between the columns specified before the IN keyword and a subquery result.
- The **IN** clause scan all records fetched from the given subquery column.
- For this query, for each **Sailors** tuple, it checks the qualification by computing the nested subquery, if **at least one tuple is** in the result of nested query, then select that **Sailors** tuple.
- To find “*sailors who’ve not reserved #103,*” use **NOT IN**.

# Nested Queries with Correlation

*Q: “Find names of sailors who’ve reserved boat #103”*

```
SELECT S.sname
FROM   Sailors S
WHERE  EXISTS (SELECT *
               FROM   Reserves R
               WHERE  R.bid=103 AND S.sid=R.sid)
```



- **EXISTS** is another set comparison operator. **EXISTS** is used to check whether the result of a correlated nested subquery is **empty** or **not**. So, it checks the subquery result and returns an either **TRUE** or **FALSE** value.
- **EXISTS** operator returns TRUE if the subquery returns single or multiple records. Otherwise, it gives a FALSE result when no records are returned.
- If **UNIQUE** is used, and \* is replaced by *R.bid*, “*finds sailors with at most one reservation for boat #103.*” (UNIQUE checks for duplicate tuples; \* denotes all attributes.)

# Examples of Division A/B

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

*A*

pno
p2

*B1*

sno
s1
s2
s3
s4

*A/B1*

pno
p2
p4

*B2*

sno
s1
s4

*A/B2*

pno
p1
p2
p4

*B3*

sno
s1

*A/B3*

# Division in SQL

Q: “Find sailors who’ve reserved *all* boats.”

- Let’s do it without EXCEPT:

(2)

SELECT S.sname     *Sailors S such that ...*

FROM     Sailors S

WHERE NOT EXISTS (SELECT B.bid     *there is no boat B*

FROM     Boats B

WHERE NOT EXISTS (SELECT R.bid

FROM     Reserves R

WHERE R.bid=B.bid

AND R.sid=S.sid))

*without any Reserves tuple showing S reserved B*

```
SELECT S.sname  
FROM   Sailors S  
WHERE  NOT EXISTS  
      ((SELECT B.bid  
        FROM   Boats B)  
EXCEPT  
      (SELECT R.bid  
        FROM   Reserves R  
        WHERE  R.sid=S.sid))
```

OR “Select each sailor such that there does not exist any boat that the sailor does not reserve it.”

# More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: *op* ANY, *op* ALL,  
$$op \text{ IN } \{ >, <, =, \geq, \leq, \neq \}$$

*Q: “Find sailors whose rating is greater than that of some sailor called Horatio”*

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating > ANY (SELECT S2.rating  
                        FROM   Sailors S2  
                        WHERE  S2.sname='Horatio')
```

# Aggregate Operators

- Significant extension of relational algebra.

```
SELECT COUNT (*)  
FROM Sailors S
```

```
SELECT S.sname  
FROM Sailors S
```

```
WHERE S.rating = (SELECT MAX(S2.rating)  
FROM Sailors S2)
```

```
SELECT AVG (S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT COUNT (DISTINCT S.rating)  
FROM Sailors S  
WHERE S.sname='Bob'
```

```
SELECT AVG (DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10
```

COUNT (\*)  
COUNT ( [DISTINCT] A)  
SUM ( [DISTINCT] A)  
AVG ( [DISTINCT] A)  
MAX (A)  
MIN (A)

*single column*



## AIRPORT

<u>airportcode</u>	name	city	state
--------------------	------	------	-------

## FLT-SCHEDULE

<u>flt#</u>	airline	dtime	from-airportcode	atime	to-airportcode	miles	price
-------------	---------	-------	------------------	-------	----------------	-------	-------

## FLT-WEEKDAY

<u>flt#</u>	<u>weekday</u>
-------------	----------------

## FLT-INSTANCE

<u>flt#</u>	<u>date</u>	plane#	#avail-seats
-------------	-------------	--------	--------------

## AIRPLANE

<u>plane#</u>	plane-type	total-#seats
---------------	------------	--------------

## CUSTOMER

<u>cust#</u>	first	middle	last	phone#	street	city	state	zip
--------------	-------	--------	------	--------	--------	------	-------	-----

## RESERVATION

<u>flt#</u>	date	<u>cust#</u>	seat#	check-in-status	<u>ticket#</u>
-------------	------	--------------	-------	-----------------	----------------

# Interactive DML - from clause

**Q:** “Find FLT#, WEEKDAY, and FROM-AIRPORTCODE in FLT-WEEKDAY and FLT-SCHEDULE”

```
SELECT FLT-SCHEDULE.FLT#, WEEKDAY, FROM-AIRPORTCODE  
FROM FLT-WEEKDAY FW, FLT-SCHEDULE FS  
WHERE FW.FLT# = FS.FLT#;
```

- **dot-notation** disambiguates FLT# in FLT-WEEKDAY and FLT-SCHEDULE
- this is a **natural join**:

$\pi$  (FLT-SCHEDULE  $\bowtie$  FLT-WEEKDAY)  
FLT-SCHEDULE.FLT#, WEEKDAY, FROM-AIRPORTCODE

# Interactive DML - ordering of rows


- the **order by** clause orders the rows in a query result in ascending (**asc**) or descending (**desc**) order

*Q: “Find FLT#, airline, and price from FLT-SCHEDULE for flights out of Atlanta ordered by ascending airline and descending price:”*

```
SELECT FLT#, AIRLINE, PRICE  
FROM FLT-SCHEDULE  
WHERE FROM-AIRPORTCODE=“ATL”  
ORDER BY AIRLINE ASC, PRICE DESC;
```

# Interactive DML - set operations

$S \cup T$



$S \text{ union } T$

*Q: “Find FLT# for flights on Tuesdays in FLT-WEEKDAY or FLT# with more than 100 seats in FLT-INSTANCE ”*


```
SELECT FLT#  
FROM FLT-WEEKDAY  
WHERE WEEKDAY = “TU”  
UNION  
SELECT FLT#  
FROM FLT-INSTANCE  
WHERE #AVAIL-SEATS > 100;
```

- **UNION**: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).

- **UNION ALL** preserves duplicates

# Interactive DML - set operation

$S \cap T$



S intersect T


**Q:** “Find FLT# for flights on Tuesdays in FLT-WEEKDAY with more than 100 seats in FLT-INSTANCE”

```
SELECT FLT#  
FROM FLT-WEEKDAY  
WHERE WEEKDAY = “TU”  
INTERSECT  
SELECT FLT#  
FROM FLT-INSTANCE  
WHERE #AVAIL-SEATS > 100;
```

- **INTERSECT ALL** preserves duplicates

- **INTERSECT**: Can be used to compute the intersection of any two *union-compatible* sets of tuples.
- Included in the SQL/92 standard, but some systems don't support it.

# Interactive DML - set operation



- “Find FLT# for flights on Tuesdays in FLT-WEEKDAY except FLT# with more than 100 seats in FLT-INSTANCE”

```
SELECT FLT#  
FROM FLT-WEEKDAY  
WHERE WEEKDAY = “TU”  
EXCEPT  
SELECT FLT#  
FROM FLT-INSTANCE  
WHERE #AVAIL-SEATS > 100;
```

- **EXCEPT ALL** preserves duplicates

# Interactive DML - nested subqueries

- Set membership: **IN**, **NOT IN**

**Q:** *“Find airlines from FLT-SCHEDULE where FLT# is not in the set of FLT#’s for flights on Tuesdays from FLT-WEEKDAY”*

```
SELECT DISTINCT AIRLINE  
FROM FLT-SCHEDULE  
WHERE FLT# NOT IN  
    (SELECT FLT#  
     FROM FLT-WEEKDAY  
     WHERE WEEKDAY = “TU”);
```

**Q:** *“Find FLT#’s for flights on Tuesdays or Thursdays from FLT-WEEKDAY”*

```
SELECT DISTINCT FLT#  
FROM FLT-WEEKDAY  
WHERE WEEKDAY IN (“TU”, “TH”);
```

# Interactive DML - nested subqueries

*Q: “Find FLT# for flights from Atlanta to Chicago with a price so low that there does not exist any cheaper flights from Birmingham to Chicago”*

```
SELECT S.FLT#  
FROM FLT-SCHEDULE S  
WHERE S.FROM-AIRPORTCODE=“ATL” AND  
      S.TO-AIRPORTCODE=“CHI” AND  
NOT EXISTS  
      (SELECT T.FLT#  
        FROM FLT-SCHEDULE T  
        WHERE T.FROM-AIRPORTCODE=“BIR” AND  
              T.TO-AIRPORTCODE=“CHI” AND  
              T.PRICE < S.PRICE);
```



# Interactive DML - nested subqueries

**Q:** *“Find FLT# for flights from Atlanta to Chicago with a price that is lower than all flights from Birmingham to Chicago”*

```
SELECT FLT#  
FROM FLT-SCHEDULE  
WHERE FROM-AIRPORTCODE=“ATL”  
      AND TO-AIRPORTCODE=“CHI”  
      AND PRICE < ALL (SELECT PRICE  
                        FROM FLT-SCHEDULE  
                        WHERE FROM-AIRPORTCODE=“BIR”  
                        AND TO-AIRPORTCODE=“CHI”);
```

# Interactive DML - joins

- **cross join:** Cartesian product
- **[inner] join:** only keeps rows that satisfy the join condition
- **left outer join:** keeps all rows from left table; fills in nulls as needed
- **right outer join:** keeps all rows from right table; fills in nulls as needed
- **full outer join:** keeps all rows from both tables; fills in nulls as needed
- **natural or on-condition** must be specified for all inner and outer joins
- **natural:** equi-join on columns with same name; one column preserved

# Interactive DML - joins

**Q:** *“Find all two-leg, one-day trips out of Atlanta; show also a leg-one even if there is no connecting leg-two the same day”*

```
SELECT X.FLT# LEG-ONE, Y.FLT# LEG-TWO
FROM ((FLT-SCHEDULE NATURAL JOIN FLT-INSTANCE) X
LEFT OUTER JOIN
  (FLT-SCHEDULE NATURAL JOIN FLT-INSTANCE) Y
ON (X.TO-AIRPORTCODE = Y.FROM-AIRPORTCODE AND
  X.DATE = Y.DATE AND X.ATIME < Y.DTIME))
WHERE X.FROM-AIRPORTCODE=“ATL”;
```

# Queries With GROUP BY and HAVING

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

- The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
  - The **attribute list** (i) must be a subset of *grouping-list*. Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a single value per group.
  - A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.

# Queries With GROUP BY and HAVING

## Conceptual Evaluation:

- The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded, 'unnecessary' fields are deleted (projected out), and the remaining tuples are partitioned into *groups* by the value of attributes in *grouping-list*.
- The *group-qualification* (having clause) is then applied to eliminate some groups.
- In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*.

**Q:** “Find the age of the youngest sailor with age  $\geq 18$ , for each rating with at least 2 such sailors”

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

- Only S.rating and S.age are mentioned in the SELECT, GROUP BY or HAVING clauses; other attributes ‘unnecessary’.
- 2nd column of result is unnamed.  
(Use AS to name it.)

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0


rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

rating	
7	35.0


*Answer relation*

**Q:** Find age of the youngest sailor with age  $\geq 18$ , for each rating with at least 2 such sailors.

rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
10	16.0
9	35.0
3	25.5
3	63.5
3	25.5



rating	age
1	33.0
3	25.5
3	63.5
3	25.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0



rating	minage
3	25.5
7	35.0
8	25.5

## Interactive DML - built-in functions

**Q:** *“Find the average ticket price by airline for scheduled flights out of Atlanta for airlines with more than 5 scheduled flights out of Atlanta from FLT-SCHEDULE”*

```
SELECT AIRLINE, AVG(PRICE)
FROM FLT-SCHEDULE
WHERE FROM-AIRPORTCODE = “ATL”
GROUP BY AIRLINE
HAVING COUNT (FLT#) >= 5;
```



# Interactive DML - insert, delete, update

```
INSERT INTO FLT-SCHEDULE  
VALUES (“DL212”, “DELTA”, 11-15-00, “ATL”,  
13-05-00, ”CHI”, 650, 00351.00);
```

```
INSERT INTO FLT-SCHEDULE(FLT#,AIRLINE)  
VALUES (“DL212”, “DELTA”); /*default nulls added*/
```

**Q:** *“Insert into FLT-INSTANCE all flights scheduled for Thursday, 9/10/98”*

```
INSERT INTO FLT-INSTANCE(FLT#, DATE)  
(SELECT S.FLT#, 1998-09-10  
FROM FLT-SCHEDULE S, FLT-WEEKDAY D  
WHERE S.FLT#=D.FLT#  
AND D.WEEKDAY=“TH”);
```

# Interactive DML - insert, delete, update

**Q:** *“Cancel all flight instances for Delta on 9/10/98”*

**DELETE FROM FLT-INSTANCE**

**WHERE DATE=1998-09-10**

**AND FLT# IN**

**(SELECT FLT#**

**FROM FLT-SCHEDULE**

**WHERE AIRLINE=“DELTA”);**

# Interactive DML - insert, delete, update

**Q:** *“Update all reservations for customers on DL212 on 9/10/98 to reservations on AA121 on 9/10/98”*

**UPDATE RESERVATION**

**SET** FLT#="AA121"

**WHERE** DATE=1998-09-10

**AND** FLT#="DL212";