

Ceng 302
Database Systems

SQL: Structured Query Language-3

Prof. Dr. Adnan YAZICI
Department of Computer Engineering,
METU
(Fall 2021)

Queries With GROUP BY and HAVING

recap

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

- The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
 - The **attribute list** (i) must be a subset of *grouping-list*. Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a single value per group.
 - A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.

Queries With GROUP BY and HAVING

recap

Conceptual Evaluation:

- The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded, 'unnecessary' fields are deleted, and the remaining tuples are partitioned into *groups* by the value of attributes in *grouping-list*.
- The *group-qualification* is then applied to eliminate some groups.
- In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*.

Q: “Find the age of the youngest sailor with age ≥ 18 , for each rating with at least 2 such sailors”

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

- Only S.rating and S.age are mentioned in the SELECT, GROUP BY or HAVING clauses; other attributes ‘unnecessary’.
- 2nd column of result is unnamed.
(Use AS to name it.)

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0


rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

rating	
7	35.0


Answer relation

Q: Find age of the youngest sailor with age ≥ 18 ,
for each rating with at least 2 such sailors.

rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
10	16.0
9	35.0
3	25.5
3	63.5
3	25.5



rating	age
1	33.0
3	25.5
3	63.5
3	25.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0



rating	minage
3	25.5
7	35.0
8	25.5

Interactive DML - built-in functions

recap

Q: *“Find the average ticket price by airline for scheduled flights out of Atlanta for airlines with more than 5 scheduled flights out of Atlanta from FLT-SCHEDULE”*

```
SELECT AIRLINE, AVG(PRICE)
FROM FLT-SCHEDULE
WHERE FROM-AIRPORTCODE = “ATL”
GROUP BY AIRLINE
HAVING COUNT (FLT#) >= 5;
```

Interactive DML - insert, delete, update

recap

```
INSERT INTO FLT-SCHEDULE  
VALUES (“DL212”, “DELTA”, 11-15-00, “ATL”,  
13-05-00, ”CHI”, 650, 00351.00);
```

```
INSERT INTO FLT-SCHEDULE(FLT#,AIRLINE)  
VALUES (“DL212”, “DELTA”); /*default nulls added*/
```

Q: *“Insert into FLT-INSTANCE all flights scheduled for Thursday, 9/10/98”*

```
INSERT INTO FLT-INSTANCE(FLT#, DATE)  
(SELECT S.FLT#, 1998-09-10  
FROM FLT-SCHEDULE S, FLT-WEEKDAY D  
WHERE S.FLT#=D.FLT#  
AND D.WEEKDAY=“TH”);
```

Interactive DML - insert, delete, update

recap

Q: “Cancel all flight instances for Delta on 9/10/98”

```
DELETE FROM FLT-INSTANCE  
WHERE DATE=1998-09-10  
AND FLT# IN  
    (SELECT FLT#  
        FROM FLT-SCHEDULE  
        WHERE AIRLINE=“DELTA”);
```


Interactive DML - insert, delete, update

recap

Q: “Update all reservations for customers on DL212 on 9/10/98 to reservations on AA121 on 9/10/98”

UPDATE RESERVATION

SET FLT#="AA121"

WHERE DATE=1998-09-10

AND FLT#="DL212";

Embedded DML- impedance mismatch

- SQL is a powerful, set-oriented, declarative language
- SQL queries return sets of rows
- Host languages cannot handle large sets of structured data
- Cursors resolve the mismatch:

EXEC SQL BEGIN DECLARE SECTION;

.....

EXEC SQL END DECLARE SECTION;

EXEC SQL

DECLARE <key attribute-name> **CURSOR**

FOR <SQL QUERY>

EXEC SQL OPEN <ATTR>;

*/*query executed; cursor open;*/*

WHILE MORE <ATTR> **DO**

*/*first row is the current row */*

EXEC SQL FETCH <ATTR>

*/*one row of query placed in */*

INTO : <Host variables>

*/*host variables*/*

<DO YOUR THING WITH THE DATA>;

END-WHILE;

EXEC SQL CLOSE <ATTR>;

Embedded DML- database access

- to access the result of the query, one row at a time, the following is used:

```
EXEC SQL BEGIN DECLARE SECTION;
DECLARE FROM-AIRPORTCODE CHAR(3);           /*input to query*/
.....
DECLARE FLT# CHAR(5);                         /*targets for FETCH*/
DECLARE AIRLINE VARCHAR(25);
DECLARE PRICE DECIMAL(7,2);
EXEC SQL END DECLARE SECTION;
EXEC SQL
DECLARE FLT CURSOR
FOR SELECT FLT#, AIRLINE, PRICE           /*SQL query;*/
           FROM FLT-SCHEDULE
           WHERE FROM-AIRPORTCODE = FROM-AIRPORTCODE;
EXEC SQL OPEN FLT;                          /*query executed; cursor open;*/
WHILE MORE FLTs DO                          /*first row is the current row */
                                           /*one row of query placed in */
           INEXEC SQL FETCH TO :FLT#, :AIRLINE, :PRICE;    /*host variables*/
           <DO YOUR THING WITH THE DATA>;
END-WHILE;
EXEC SQL CLOSE FLT;                        /*cursor closed*/
```

Views- definition, use, update

What is a view?

- A view is a virtual table.

```
Create View Vname As  
<Query>
```

Why use views?

- Simplify Queries
- Hide some data from some users (security)
- Make some queries easier / more natural
- Modularity of database access
- Improve performance of database queries (only if materialized views used)

Querying views

- Once *View* is defined, one can reference *View* like any table
- Queries involving *View* rewritten to use base tables

Types of Views

- **Virtual views:** ✓
 - Used in databases ✓
 - Computed only on-demand – slower at runtime
 - Always up to date ✓
- **Materialized views** ✓
 - Used in data warehouses, OLAP,
 - Precomputed offline – faster at runtime
 - May have stale data

Views- definition, use, update

- how a view is defined:

```
CREATE VIEW ATL-FLT  
AS SELECT FLT#, AIRLINE, PRICE  
FROM FLT-SCHEDULE  
WHERE FROM-AIRPORTCODE = “ATL”;
```

- how a query on a view is written:

```
SELECT *  
FROM ATL-FLT  
WHERE PRICE <= 00200.00;
```

- how this query on a view is computed:

```
SELECT FLT#, AIRLINE, PRICE  
FROM FLT-SCHEDULE  
WHERE FROM-AIRPORTCODE=“ATL” AND PRICE<00200.00;
```

- how a view definition is dropped:

```
DROP VIEW ATL-FLT [RESTRICT|CASCADE];
```

Views- definition, use, update

Materialized views:

- View $V = \text{ViewQuery}(R_1, R_2, \dots, R_n)$
- Create table V with schema of query result
- Execute ViewQuery and put results in V
- Queries refer to V as if it's a table

But...

- V could be very large
- Modifications to $R_1, R_2, \dots, R_n \Rightarrow$
recompute or modify V
- Some DBMS allow views to be stored (*materialised views*)
 - materialised views have to be updated when its relations change (*view maintenance*)

Materialized Views

```
Create Materialized View CA-CS As
Select C.cName, S.sName
From College C, Student S, Apply A
Where C.cName = A.cName And S.sID = A.sID
And C.state = 'CA' And A.major = 'CS'
```

- + Can use **CA-CS** as if it's a table (it is!)
- + Modifications to base data invalidate view **CA-CS**?

College

cName	state	enr

Student

sID	sName	GPA	HS

Apply

sID	cName	major	dec

Views- definition, use, update

Modifying views

- Once *View* defined, can we **modify** *View* like any other table ?
- One way of thinking of this is that it **doesn't** make sense, when *View* is not stored
- The other way of thinking of this is that it **has to** make sense when *Views* are some users' entire "view" of the database

❖ **Soln: Modifications to *View* rewritten to modify base tables**

(1) Rewriting process specified explicitly by view creator

+ Can handle all modifications

– No guarantee of correctness (or meaningful)

(2) Restrict views + modifications so that translation to base table modifications is meaningful and unambiguous

+ No user intervention

– Restrictions are significant

Views- definition, use, update

a view is **updatable if and only if**:

- it does not contain any of the keywords **JOIN, UNION, INTERSECT, EXCEPT**
- it does not contain the keyword **DISTINCT**
- every column in the view corresponds to a uniquely identifiable base table column
- Attributes not in view can be **NULL** or have default value
- the **FROM** clause references exactly one table which must be a base table or an updatable view
- the table referenced in the **FROM** clause in the view cannot be referenced in the **FROM** clause of a nested **WHERE** clause
- it does not have a **GROUP BY** clause or aggregation
- it does not have a **HAVING** clause

Updatable means insert, delete, update all ok

Views- definition, use, update

```
CREATE VIEW LOW-ATL-FARES           /* updatable view*/  
AS SELECT *  
  FROM FLT-SCHEDULE  
  WHERE FROM-AIRPORTCODE="ATL"  
  AND PRICE<00200.00;
```

```
UPDATE LOW-ATL-FARES           /* moves row      */  
SET   PRICE = 00250.00         /* outside the view*/  
WHERE TO-AIRPORTCODE = "BOS";
```

```
INSERT INTO LOW-ATL-FARES       /* creates row      */  
VALUES ("DL222", "DELTA",      /* outside the view*/  
        "BIR", 11-15-00, "CHI", 13-05-00, 00180.00);
```

```
CREATE VIEW LOW-ATL-FARES  
AS SELECT *  
  FROM FLT-SCHEDULE  
  WHERE FROM-AIRPORTCODE="ATL"  
  AND PRICE<00200.00  
WITH CHECK OPTION;           /* prevents updates*/  
                               /* outside the view*/
```

Views- definition, use, update

these views are **not** updatable

```
CREATE VIEW ATL-PRICES  
AS SELECT AIRLINE, PRICE  
FROM FLT-SCHEDULE  
WHERE FROM-AIRPORTCODE="ATL";
```

not unique



computed attribute



```
CREATE VIEW AVG-ATL-PRICES  
AS SELECT AIRLINE, AVG(PRICE)  
FROM FLT-SCHEDULE  
WHERE FROM-AIRPORTCODE="ATL"  
GROUP BY AIRLINE;
```

this view is theoretically updatable, but cannot be updated in SQL

```
CREATE VIEW FLT-SCHED-AND-DAY  
AS SELECT S.*, D.WEEKDAY  
FROM FLT-SCHEDULE S, FLT-WEEKDAY D  
WHERE D.FLT# = S.FLT#;
```

join of two tables



Views- definition, use, update

Queries over materialized views

- View $V = \text{ViewQuery}(R_1, R_2, \dots, R_n)$
- Create table V with schema of query result
- Execute ViewQuery and put results in V
- Queries refer to V as if it's a table

Modifications on materialized views?

- Good news: just update the stored table
- Bad news: base tables must stay in synchronized

❖ Same issues as with virtual views

Picking which materialized views to create

(Efficiency) benefits of a materialized view depend on:

- Size of data
 - Complexity of view
 - Number of queries using view
 - Number of modifications affecting view
-
- Also “incremental maintenance” versus full recomputation

Integrity constraints

Impose restrictions on allowable data,
beyond those imposed by structure and types

- Non-null constraints
- Key constraints
- Attribute-based and tuple-based constraints
- Referential Integrity (foreign key) constraints
- General assertions

Integrity- constraints

- **constraint**: a conditional expression required not to evaluate to *false*
- a constraint cannot be created if it is already violated
- a constraint is enforced from the point of creation forward
- a constraint has a unique name
- if a constraint is violated, its name is made available to the user
- constraints cannot reference parameters or host variables; they are application independent.
- data type checking is a primitive form of constraint

Integrity Constraints

- An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts/deletes/updates that violate IC's are disallowed.
 - Can be used to ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 200)
- *Types of IC's*: Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - *Domain constraints*: Field values must be of right type. Always enforced.

Integrity- domain constraints

- associated with a domain; applies to all columns defined on the domain

```
CREATE DOMAIN WEEKDAY CHAR(2)  
CONSTRAINT IC-WEEKDAY  
CHECK (VALUE IN  
    ( "MO", "TU", "WE", "TH", "FR", "SA", "SU" ));
```

```
CREATE DOMAIN PRICE DECIMAL(7,2)  
CONSTRAINT IC-PRICE  
CHECK (VALUE > 00000.00 );
```

```
CREATE DOMAIN FLT# CHAR(5)  
CONSTRAINT IC-FLT#  
CHECK (VALUE NOT NULL);
```

Integrity- base table, column constraints

- associated with a specific base table

```
CREATE TABLE AIRLINE.FLT-SCHEDULE
(FLT#          FLIGHTNUMBER NOT NULL,
AIRLINE        VARCHAR(25),
FROM-AIRPORTCODE  AIRPORT-CODE,
DTIME          TIME,
TO-AIRPORTCODE   AIRPORT-CODE,
ATIME          TIME,
CONSTRAINT FLTPK PRIMARY KEY (FLT#),
CONSTRAINT FROM-AIRPORTCODE-FK
FOREIGN KEY (FROM-AIRPORTCODE)
REFERENCES AIRPORT(AIRPORTCODE)
ON DELETE SET NULL ON UPDATE CASCADE,
FOREIGN KEY (FROM-AIRPORTCODE)
REFERENCES AIRPORT(AIRPORTCODE)
ON DELETE SET NULL ON UPDATE CASCADE,
CONSTRAINT IC-DTIME-ATIME
CHECK (DTIME < ATIME);
```

General Constraints

- Useful when more general ICs than keys are involved.
- Can use queries to express constraint.
- Constraints can be named.

```
CREATE TABLE Sailors
```

```
(sid INTEGER,  
sname CHAR(10),  
rating INTEGER,  
age REAL,  
PRIMARY KEY (sid),  
CHECK ( rating >= 1  
AND rating <= 10 )
```

```
CREATE TABLE Reserves
```

```
( sname CHAR(10),  
bid INTEGER,  
day DATE,  
PRIMARY KEY (bid,day),  
CONSTRAINT noInterlakeReserv  
CHECK (`Interlake' <>  
    ( SELECT B.bname  
      FROM Boats B  
      WHERE B.bid=bid)))
```

Database Authorization

- Make sure users see only the data that they are supposed to see
- Guard the database against modifications by malicious users
- Users have privileges; can only operate on data for which they are authorized
- **Grant** and **Revoke** statements
- Beyond simple table-level privileges:
use views privacy

Authorization

Obtaining Privileges

- Relation creator is owner
- Owner has all privileges and may grant or revoke privileges
- Discretionary Access Control (DAC) is supported by **GRANT** and **REVOKE**:

GRANT <privileges>

ON <table>

TO <users>

[**WITH GRANT OPTION**];

REVOKE [**GRANT OPTION FOR**] <privileges>

ON <table>

FROM <users> {**RESTRICT** | **CASCADE**};

<privileges>: SELECT, INSERT(X), INSERT, UPDATE(X), UPDATE, DELETE

<users>: public, U2, U8..

CASCADE: revoke cascades through its subtree

RESTRICT: revoke succeeds only if there is no subtree

Authorization

GRANT INSERT, DELETE
ON FLT-SCHEDULE
TO U1, U2
WITH GRANT OPTION;

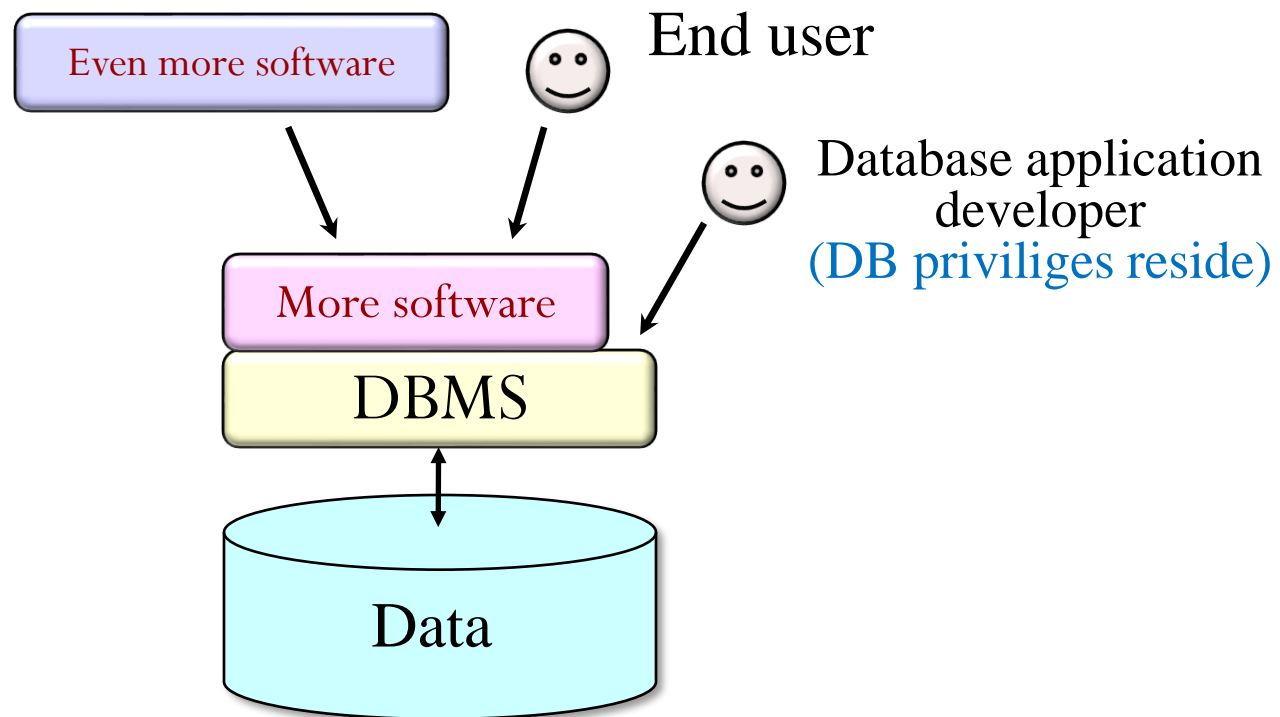
GRANT UPDATE(PRICE)
ON FLT-SCHEDULE
TO U3;

REVOKE GRANT OPTION FOR DELETE
ON FLT-SCHEDULE
FROM U2 **CASCADE;**

REVOKE DELETE
ON FLT-SCHEDULE
FROM U2 **CASCADE;**

Database Authorization

Where Privileges Reside



Summary

- SQL was an important factor in the early **acceptance** of the relational model; more natural than earlier, procedural query languages.
- **Relationally complete**; in fact, significantly more expressive power than relational algebra.
- Even queries that can be expressed in Relational Algebra can often be expressed more naturally in SQL.
- Many alternative ways to write a query; **optimizer** should look for most efficient evaluation plan.
 - In practice, users need to be aware of how queries are optimized and evaluated for best results.
- NULL for unknown field values brings many complications
- SQL allows specification of **rich integrity constraints**
- **Triggers** respond to changes in the database
- SQL supports linear recursion to increase the expressiveness.