

**Ceng 302**  
**Database Management Systems**

**Physical Data Storage Devices**

**Prof. Dr. Adnan YAZICI**

Department of Computer Engineering,  
Middle East Technical University  
(Fall 2021)

# Secondary Storage Devices

Two major types of storage devices:



1. Direct Access Storage Devices (DASDs)
  - **Magnetic Disks**
    - Hard disks (high capacity, low cost per bit)
    - Floppy disks (low capacity, slow, cheap)
  - **Optical Disks**
    - CD-ROM = (Compact disc, read-only memory)
2. Serial Devices
  - **Magnetic tapes** (very fast sequential access)

# Magnetic Disks


- Bits of data (0's and 1's) are stored on circular magnetic platters called **disks**.
- A disk rotates rapidly (& never stops).
- A **disk head** reads and writes bits of data as they pass under the head.
- Often, several platters are organized into a **disk pack** (or **disk drive**).

# Magnetic Disks

- A read/write head travels across a spinning magnetic disk, retrieving or recording data
- Each disk surface is divided into sectors and tracks
- Example of disk addressing scheme: surface 3, sector 5, track 4



# A Disk Drive



Disk drive with 4 platters and 8 surfaces

# Organization of Disks

- Disk contains concentric **tracks**.
- Tracks are divided into **sectors**
- A **sector** is the smallest addressable unit in a disk.

# Accessing Data


- When a program reads a byte from the disk, the operating system locates the surface, track and sector containing that byte, and reads the **entire sector** into a special area in main memory called **buffer**.
- The bottleneck of a disk access is moving the read/write arm. So, it makes sense to store a file in tracks that are below/above each other in different surfaces, rather than in several tracks in the same surface.

# Cylinders

- A **cylinder** is the set of tracks at a given radius of a disk pack.
  - i.e., a cylinder is the set of tracks that can be accessed without moving the disk arm.
- All the information on a cylinder can be accessed without moving the read/write arm.



# Cylinders



# Estimating Capacities

**Track capacity** = # of sectors/track \* bytes/sector

**Cylinder capacity** = # of tracks/cylinder \* track capacity

**Drive capacity** = # of cylinders \* cylinder capacity

**Number of cylinders** = # of tracks in a surface

# Exercise

- Store a file of 20000 records on a disk with the following characteristics:
    - # of bytes per sector = 512
    - # of sectors per track = 40
    - # of tracks per cylinder = 12
    - # of cylinders = 1331
- Q1.** How many cylinders does the file require if each data record requires 256 bytes?
- Q2.** What is the total capacity of the disk?

# Clusters

- Another view of sector organization is the one maintained by the O.S.'s **file manager**.
- It views the file as a series of **clusters** of sectors.
- File manager uses a **file allocation table (FAT)** to map logical sectors of the file to the physical clusters.

# Extents

- If there is a lot of room on a disk, it may be possible to make a file consist entirely of **contiguous clusters**. Then we say that the file is one **extent**. (very good for sequential processing)
- If there isn't enough contiguous space available to contain an entire file, the file is divided into two or more noncontiguous parts. Each part is an extent.

# Fragmentation

**Internal fragmentation:** loss of space within a sector or a cluster.

- 1) **Due to records not fitting exactly in a sector:**  
e.g., Sector size is 512 and record size is 300 bytes. Either
  - store one record per sector, or
  - allow records *span* sectors.
- 2) **Due to the use of clusters:** If the file size is not a multiple of the cluster size, then the last cluster will be partially used.

# Choice of cluster size

Some operating systems allow system administrator to choose cluster size.

- When to use large cluster size?
- What about small cluster size?
  - The greater the block-size, the greater potential amount of internal track fragmentation.
  - The flexibility introduced by the use of blocks rather than sectors can save time since it lets the programmer determine, to a large extent, how the data is to be organized physically on disk.

# Organizing Tracks by Block

- **Disk tracks** may be divided into **user-defined blocks** rather than into **sectors**.
- **Blocks** can be fixed or variable length.
- A **block** is usually organized to hold an integral number of logical records.
- **Blocking Factor** = the number of records stored in a block.
- No internal fragmentation, no record spanning two blocks.
- In block-addressing scheme each block of data is accompanied by one or more *subblocks* containing extra information about the block.



# Non-data Overhead

- Both blocks and sectors require non-data overhead (written during formatting)
- On **sector addressable disks** this information involves sector address, track address, and condition (usable/defective). Also pre-formatting involves placing gaps and synchronization marks.
- On **block-organized disk**, more information is needed and the programmer should be aware of some of this information.

# Non-data Overhead

- Whether using a **block** or a **sector** organization, some space on the disk is taken up by non-data overhead. i.e., information stored on the disk during pre-formatting.
- On **Sector-Addressable disks**, pre-formatting involves storing, at the beginning of each sector, sector address, track address and condition (usable or defective) + gaps and synchronization marks between fields of info to help the read/write mechanism distinguish between them.
- On **Block-Organized disks**, subblock + interblock gaps have to be provided with every block.
- The relative amount of non-data space necessary for a block scheme is higher than for a sector-scheme.

# The Cost of a Disk Access

- The time to access a sector in a track on a surface is divided into 3 components:

<i><b>Time Component</b></i>	<i><b>Action</b></i>
<b>Seek Time</b>	Time to move the read/write arm to the correct cylinder
<b>Rotational delay (or latency)</b>	Time it takes for the disk to rotate so that the desired sector is under the read/write head
<b>Transfer time</b>	Once the read/write head is positioned over the data, this is the time it takes for transferring data

# Seek time

- **Seek time** is the time required to move the arm to the correct cylinder.
- Largest in cost.

## Typically:

- 5 ms (milliseconds) to move from one track to the next (track-to-track)
- 50 ms maximum (from inside track to outside track)
- 30 ms average (from one random track to another random track)

# Latency (rotational delay)

- **Latency** is the time needed for the disk to rotate so the sector we want is under the read/write head.
- Hard disks usually rotate at about 5000 rpm, which is one revolution per 5 msec.
- **Note:**
  - Min latency = 0
  - Max latency = Time for one disk revolution
  - **Average latency ( $r$ )** =  $(\text{min} + \text{max}) / 2$   
=  $\text{max} / 2$   
= time for  $\frac{1}{2}$  disk revolution
- Typically, 2 – 3 ms average

# Transfer Time

- **Transfer time** is the time for the read/write head to pass over a block.
- The transfer time is given by the formula:

$$\text{Transfer time} = \frac{\text{number of bytes transferred}}{\text{number of bytes on a track}} \times \text{rotation time}$$

- e.g., if there are 63 sectors per track, the time to transfer one sector would be 1/63 of a revolution.

# Sequential and Random Reading

Reading  $b$  blocks:

i. Sequentially:

$$\underbrace{s + r + b * \text{btt}}$$

insignificant for large files

$$\Rightarrow b * \text{btt}$$

ii. Randomly:

$$b * (s + r + \text{btt})$$

**Block Transfer Time (btt):** Time needed to read block to memory (buffer).

**Example:**  $t$  (**data transfer speed** / **data rate**) = 3000 bytes/msec

$b$ : denote the number of blocks to be read

# Sequential Reading

- Given the following disk:
  - $s = 16 \text{ ms}$
  - $r = 8.3 \text{ ms}$
  - Block transfer time =  $8.4 \text{ ms}$
- a) Calculate the time to read 10 sequential blocks
- b) Calculate the time to read 100 sequential blocks



# Random Reading

Given the same disk,

- a) Calculate the time to read 10 blocks randomly
- b) Calculate the time to read 100 blocks randomly