

Ceng 302
Database Management Systems

Database B+ Tree Index Structures


Prof. Dr. Adnan YAZICI
Department of Computer Engineering,
Middle East Technical University
(Fall 2021)

Search Trees and B Trees

- **B-tree** is one of the most important data structures in computer science.
- What does B stand for? (Not binary!)
- B-tree is a multiway search tree.
- Several versions of B-trees have been proposed, but only **B+-Trees** has been used with large files.
- A B+-tree is a B-tree in which data records are in leaf nodes, and faster sequential access is possible.

Search Trees and B Trees

- Search tree used to guide search for a record
 - Given value of one of record's fields




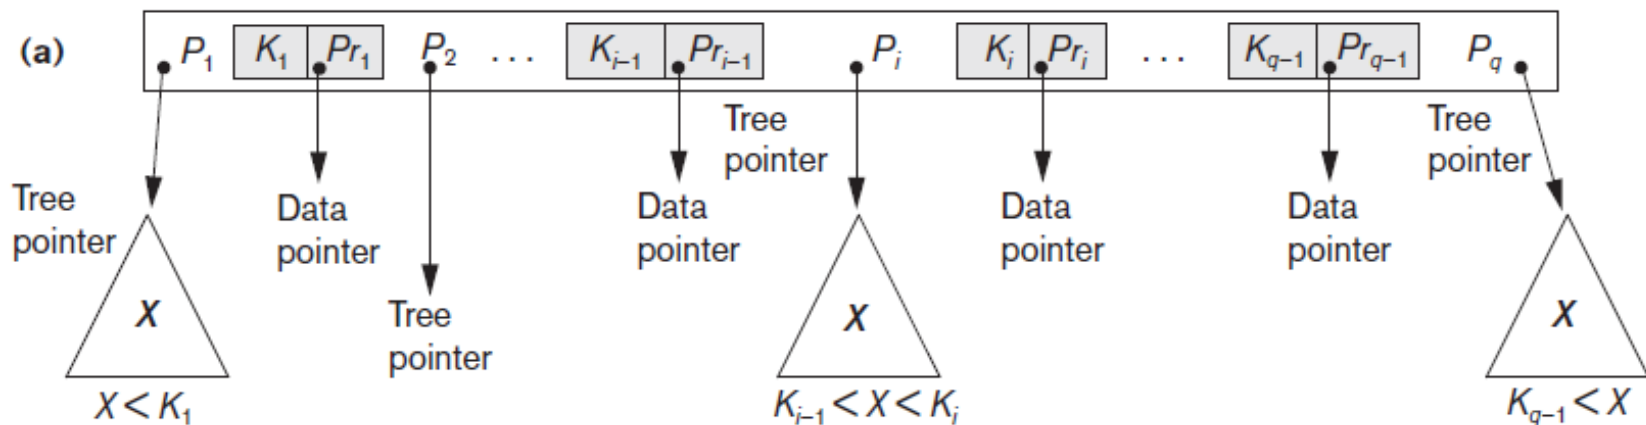
This restriction can be relaxed. If the index is on a nonkey field, duplicate search values may exist and the node structure and the navigation rules for the tree may be modified.

A node in a search tree with pointers to subtrees below it

B Trees

- Provide multi-level access structure
- Tree is always balanced
- Space wasted by deletion never becomes excessive
 - Each node is at least half-full
- Each node in a B-tree of **order p** can have at most **$p-1$ search values** and **p pointers**.

B Tree Structures



B-tree structures (a) A node in a B-tree with $q-1$ search values (b) A B-tree of **order $p=3$** . The values were inserted in the order 8, 5, 1, 7, 3, 12, 9, 6

Formal definition of B+ Tree Properties

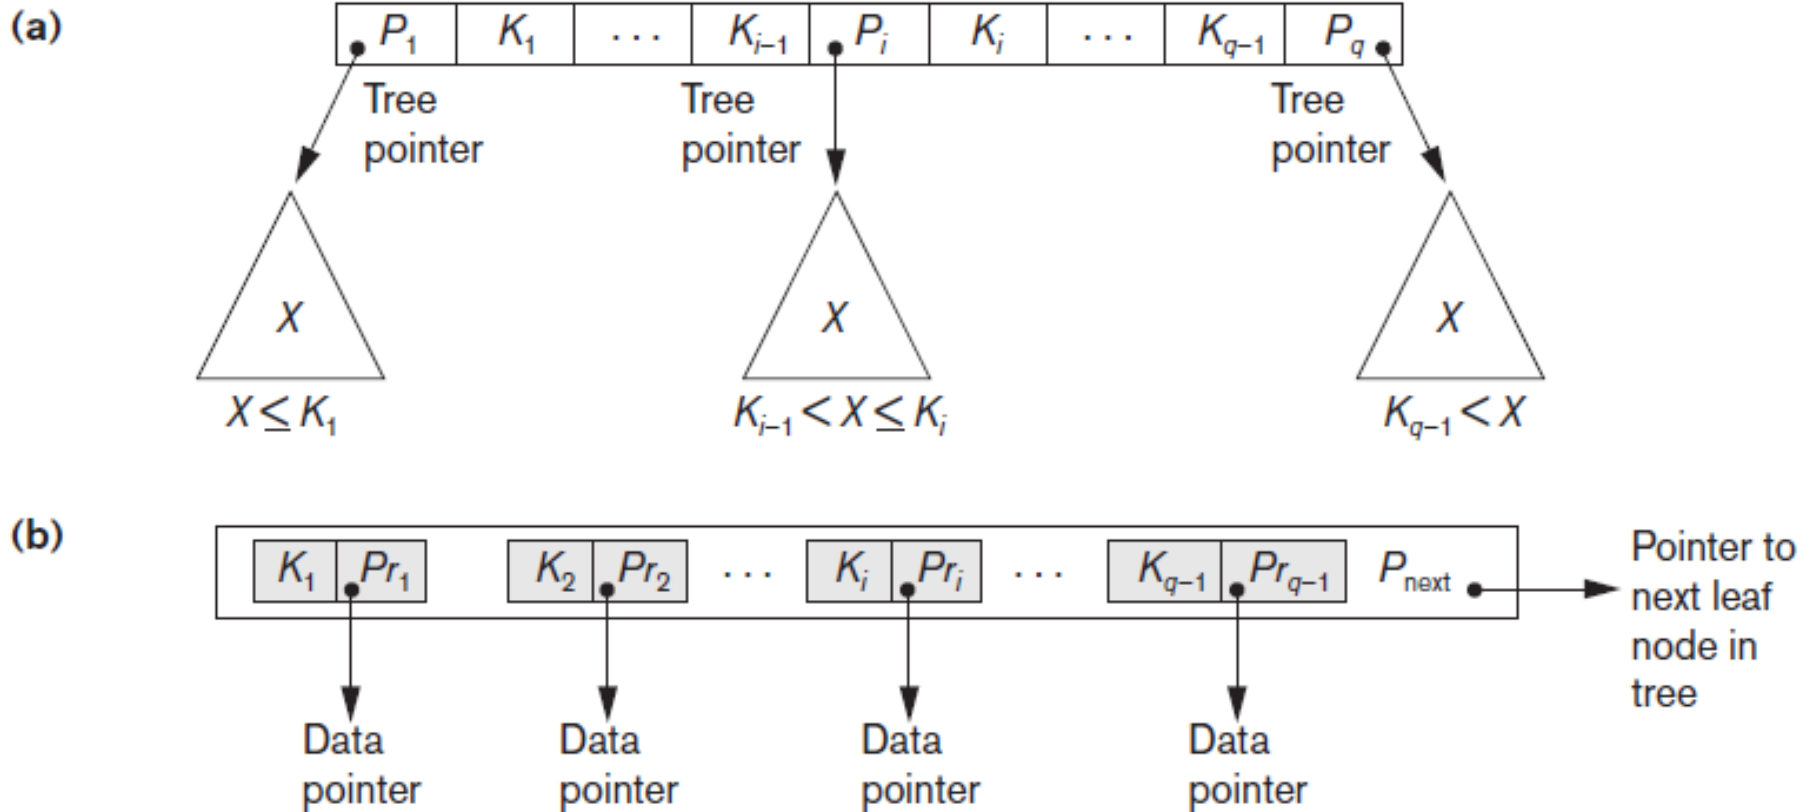
Properties of a B+ Tree of order $n+1$:

- All internal nodes (except root) has at least $n/2$ keys and at most n keys. That is, if the **order** of a B+-tree is **$n+1$** , each node (except for the root) must have between **$n/2$** and **n** keys.
- The **root** has at least **2** children unless it's a leaf.
- All **leaves** are on **the same** level.
- An **internal node** with **n keys** has **$n+1$ pointers or children**.

Terminology


- **Bucket Factor:** the number of records which can fit in a leaf node.
- **Fan-out:** the average number of children of an internal node.
- A **B+-tree index** can be used either as a **primary** index or a **secondary** index.
 - **Primary index:** determines the way the records are actually stored (also called a sparse index)
 - **Secondary index:** the records in the file are not grouped in buckets according to keys of secondary indexes (also called a dense index)

B+ Trees (cont'd.)




The nodes of a B+-tree (a) Internal node of a B+-tree with $q-1$ search values (b) Leaf node of a B+-tree with $q-1$ search values and $q-1$ data pointers

B+ Tree Indexes

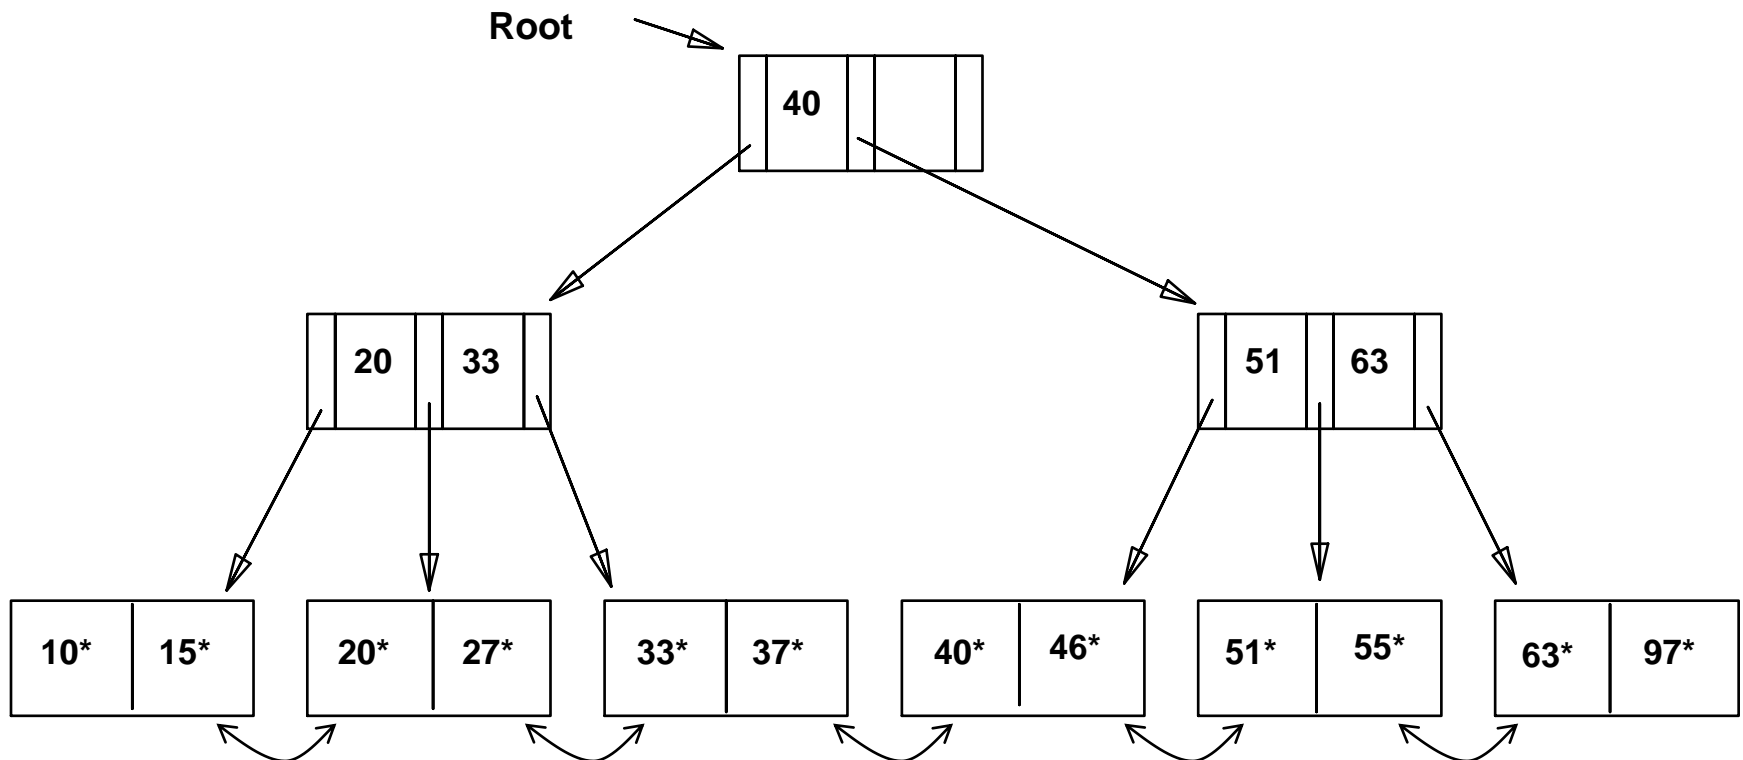


- ❖ Index leaf pages contain *data entries*, and are chained (prev & next)
- ❖ Index non-leaf pages have *index entries*; only used to direct searches:




Example: B+ tree with order of 3

- Each node must hold at least 1 entry, and at most 2 entries, and 3 pointers.




Example: Search in a B+ tree order 5


- **Search:** how to find the records with a given search key value?
 - Begin at root, and use key comparisons to go to leaf
- **Examples:** search for 5*, 16*, all data entries $\geq 24^*$...
 - The last one is a **range search**, we need to do the sequential scan, starting from the first leaf containing a value ≥ 24 .



Inserting 16*, 8* into Example B+ tree




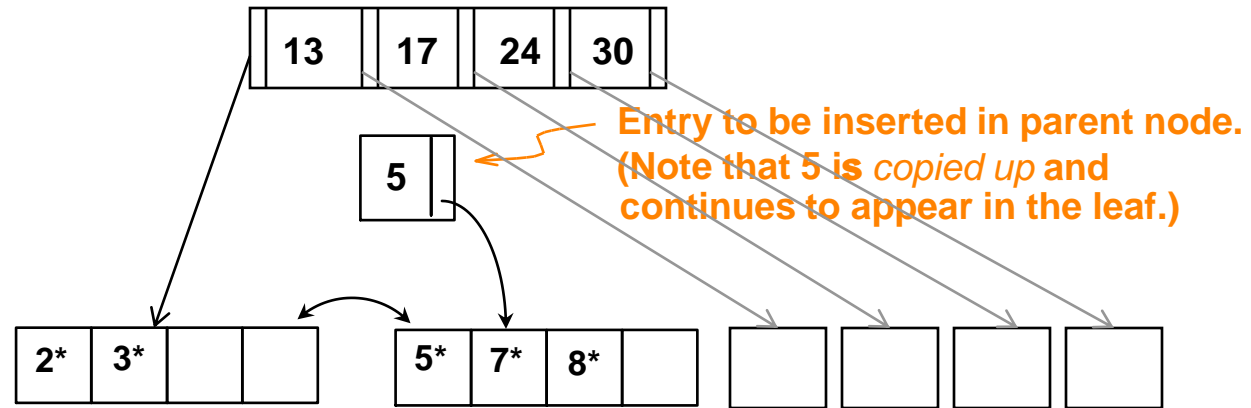
You overflow



One new child (leaf node) generated; must add one more pointer to its parent, thus one more key value as well


Inserting 8* (cont.)

- **Copy up** the middle value (leaf split)




Insertion into B+ tree (cont.)


- Understand difference between **copy-up** and **push-up**
- Observe how minimum occupancy is guaranteed in both leaf and index pg splits.



We split this node, redistribute entries evenly, and **push up** middle key.



Example B+ Tree After Inserting 8*

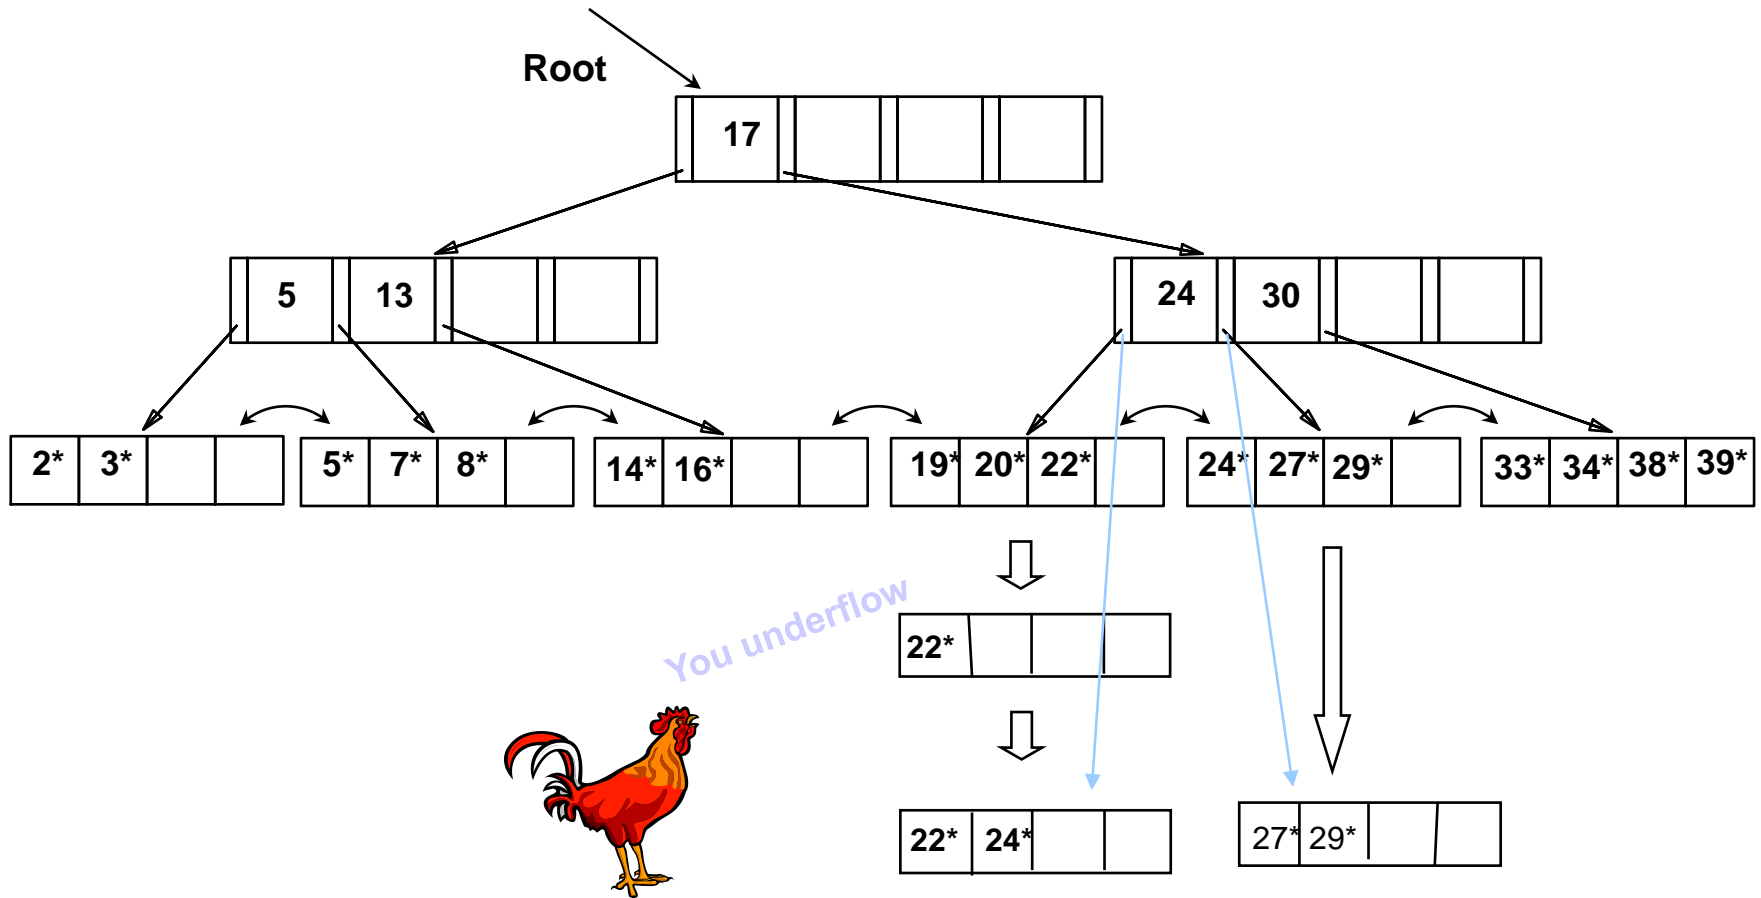


Notice that root was split, leading to increase in height.

Inserting a Data Entry into a B+ Tree: Summary


- Find correct leaf L .
- Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must *split* L (into L and a new node $L2$)
 - Redistribute entries evenly, put middle key in $L2$
 - **copy up** middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- This can happen recursively
 - To **split index node**, redistribute entries evenly, but **push up** middle key. (Contrast with leaf splits.)
- Splits “grow” tree; root split increases height.
 - Tree growth: gets *wider* or *one level taller at top*.

Delete 19* and 20*




Have we still forgot something?
CEng302-Indexing

Deleting 19* and 20* (cont.)





- Notice how 27 is *copied up*.
- But can we move it up?
- Now we want to delete 24
- Underflow again! But can we redistribute this time?

Deleting 24*



- Observe the two leaf nodes are merged, and 27 is discarded from their parent, but ...
- Observe *'pull down'* of index entry (below).




Deleting a Data Entry from a B+ Tree: Summary

- Start at root, find leaf L where entry belongs.
- Remove the entry.
 - If L is at least half-full, *done!*
 - If L has only $(n/2)-1$ entries,
 - Try to **re-distribute**, borrowing from sibling (*adjacent node with same parent as L*).
 - If re-distribution fails, **merge** L and sibling.
- If merge occurred, must delete entry (pointing to L or sibling) from parent of L .
- Merge could propagate to root, decreasing height.


Example of Non-leaf Re-distribution

- Tree is shown below *during deletion* of 24*. (What could be a possible initial tree?)
- In contrast to previous example, can re-distribute entry from left child of root to right child.



After Re-distribution

- Intuitively, entries are **re-distributed by *'pushing through'*** the splitting entry in the parent node.
- It suffices to re-distribute index entry with key 20; we've re-distributed 17 as well for illustration.



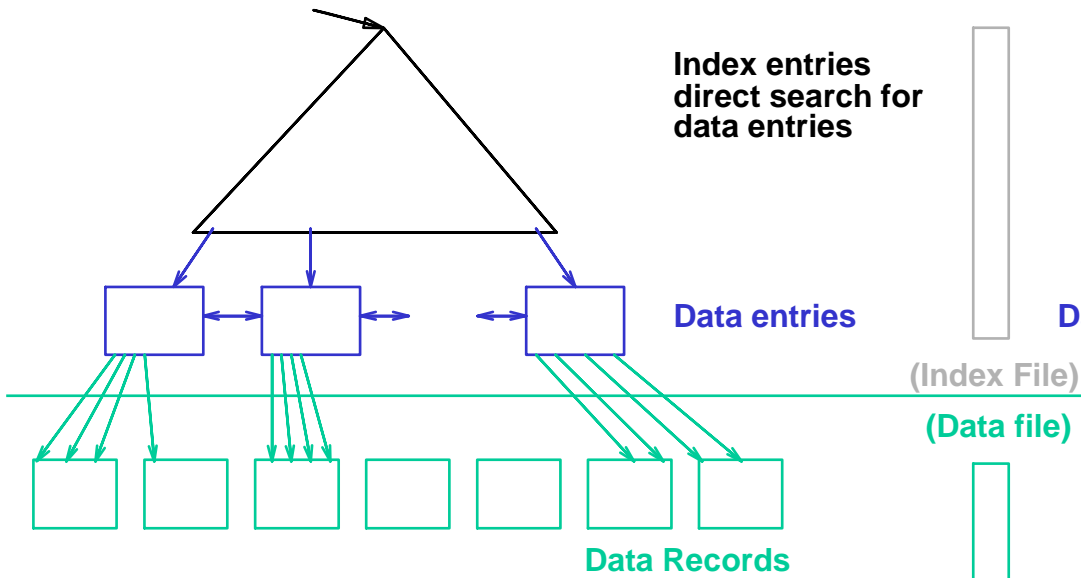
B+-Tree Animation

Index Types


- **Primary and Clustered index**
 - Controls the physical order of rows
 - Does not require disk space
 - One per table (may inc. multiple columns)
 - Created by default on tables' Primary Key column
- **Secordary (Unclustered) Index**
 - Physical data structures that facilitate data retrieval
 - Can have many indexes
 - Indexes may include many columns

Primary and Clustered vs. Secondary (Unclustered) Index

Primary or clustered index



Secondary (Unclustered) index

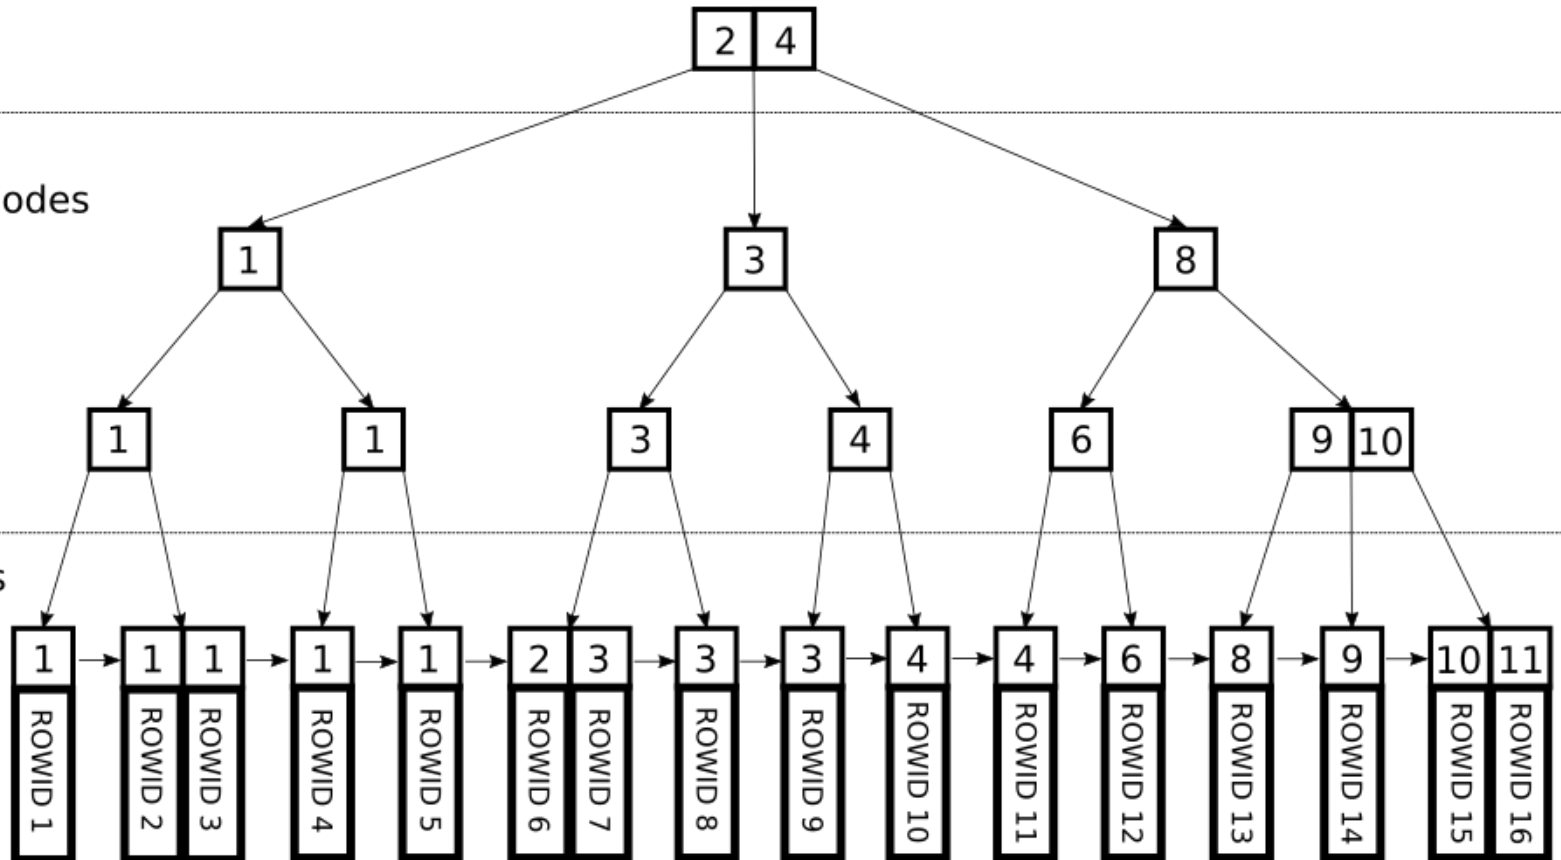


B+ Trees (cont'd.)

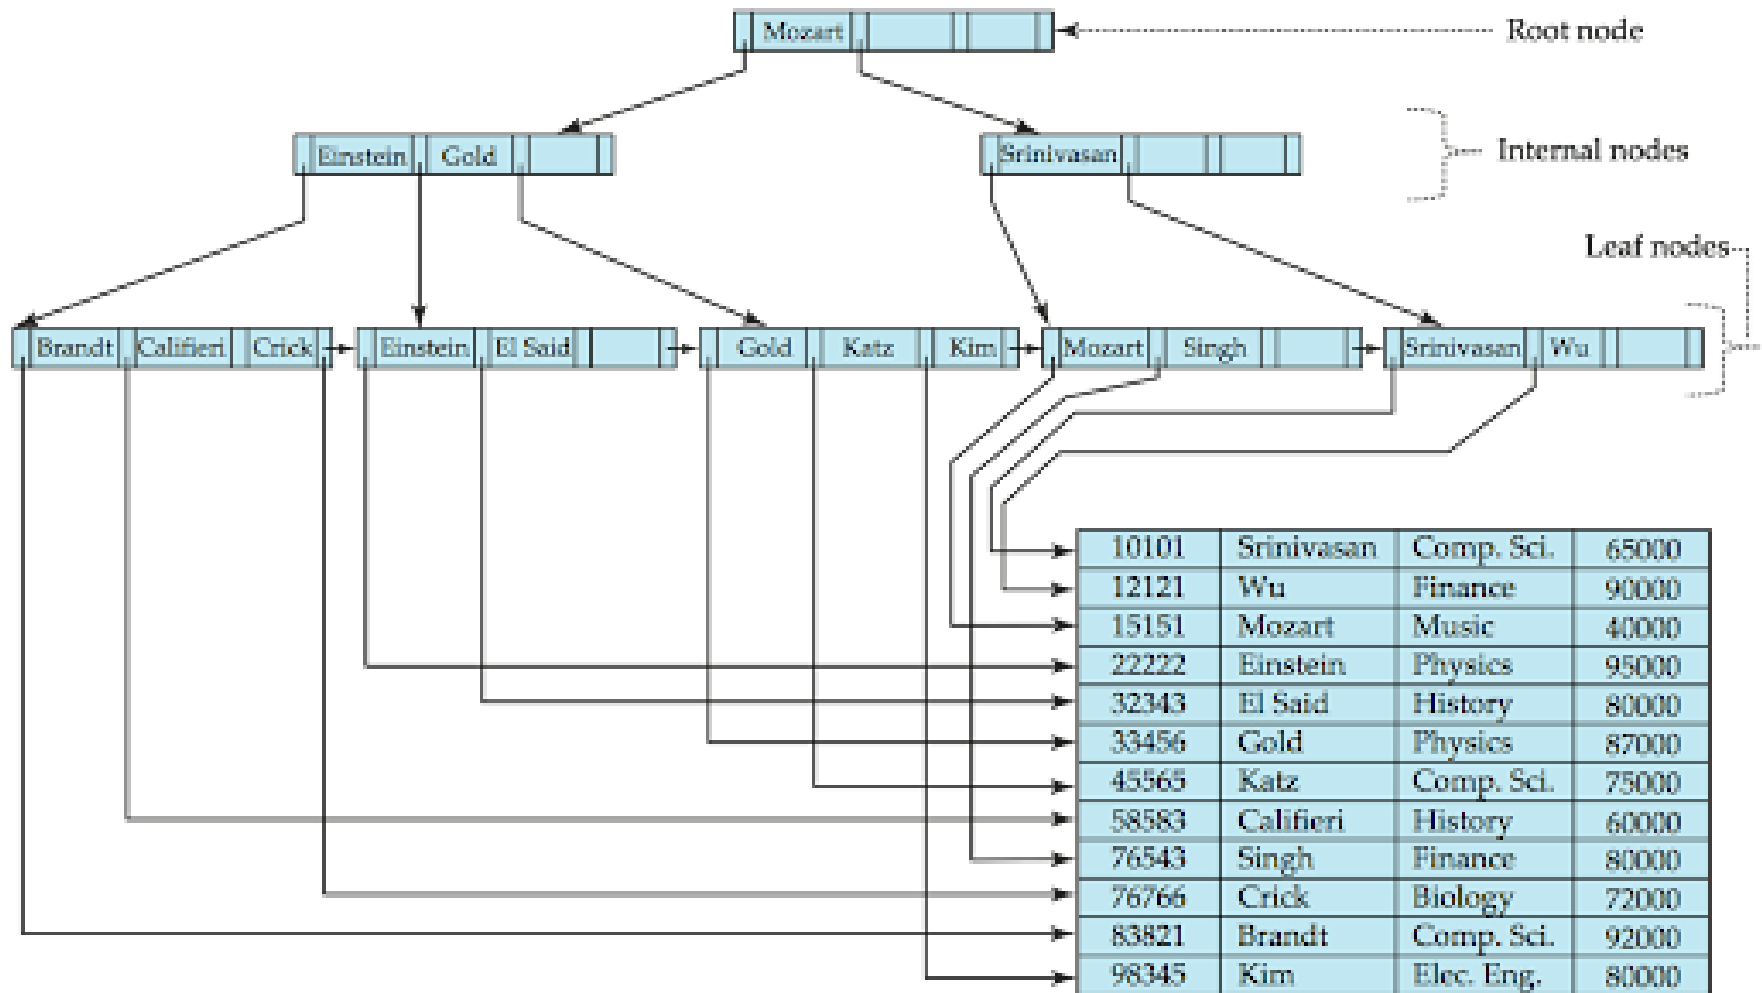
root node

non-leaf nodes

leaf nodes



B+ Trees (cont'd.)



B+ Trees (cont'd.)

- Properties
 - maximum branching factor of M
 - the root has between 2 and M children *or* at most L keys/values
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only search keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys/values
 - all leaves are at the same depth
- Result
 - height is $\Theta(\log_M n)$ between $\log_{M/2} (2n/L)$ and $\log_M (n/L)$
 - all operations run in $\Theta(\log_M n)$ time
 - operations get about $M/2$ to M or $L/2$ to L items at a time

Analysis of B+-Tree

- The maximum number of items in a B-tree of order m and height h :

root $m - 1$

level 1 $m(m - 1)$

level 2 $m^2(m - 1)$

...

level h $m^h(m - 1)$

- So, the total number of items is

$$(1 + m + m^2 + m^3 + \dots + m^h)(m - 1) =$$

$$[(m^{h+1} - 1) / (m - 1)] (m - 1) = m^{h+1} - 1 = n \Rightarrow h = \log_m n$$

- When $m = 5$ and $h = 2$ this gives $5^3 - 1 = 124$

Analysis of B+-Tree

Time complexity in big O notation		
	Average	Worst case
Space	$O(n)$	$O(n)$
Search	$O(\log n)$	$O(\log n)$
Insert	$O(\log n)$	$O(\log n)$
Delete	$O(\log n)$	$O(\log n)$

Index Creation

- General form of the command to create an index

```
CREATE [ UNIQUE ] INDEX <index name>  
ON <table name> ( <column name> [ <order> ] { , <column name> [ <order> ] } )  
[ CLUSTER ] ;
```

- Unique and cluster keywords are optional
 - Order can be ASC or DESC
- Secondary indexes can be created for any primary record organization
 - Complements other primary access methods

Definition of indexes in SQL

- The command for creating an index in SQL has the following format:

```
CREATE INDEX IndexName  
ON RelationName (ColumnNameList) [ClusterName  
[ASC | DESC];
```

where

- IndexName is the name of the index being created
- The ON clause specifies the object on which the index is allocated

Ex:

```
CREATE INDEX nameIndex ON Person(name)
```


- An index created on more than one column is called *composite index*

Indexes on Multiple Keys

- Multiple attributes involved in many retrieval and update requests
- Composite keys
 - Access structure using key value that combines attributes
- Partitioned hashing
 - Suitable for equality comparisons

Indexes on Multiple Keys (cont'd.)

- Grid files
 - Array with one dimension for each search attribute

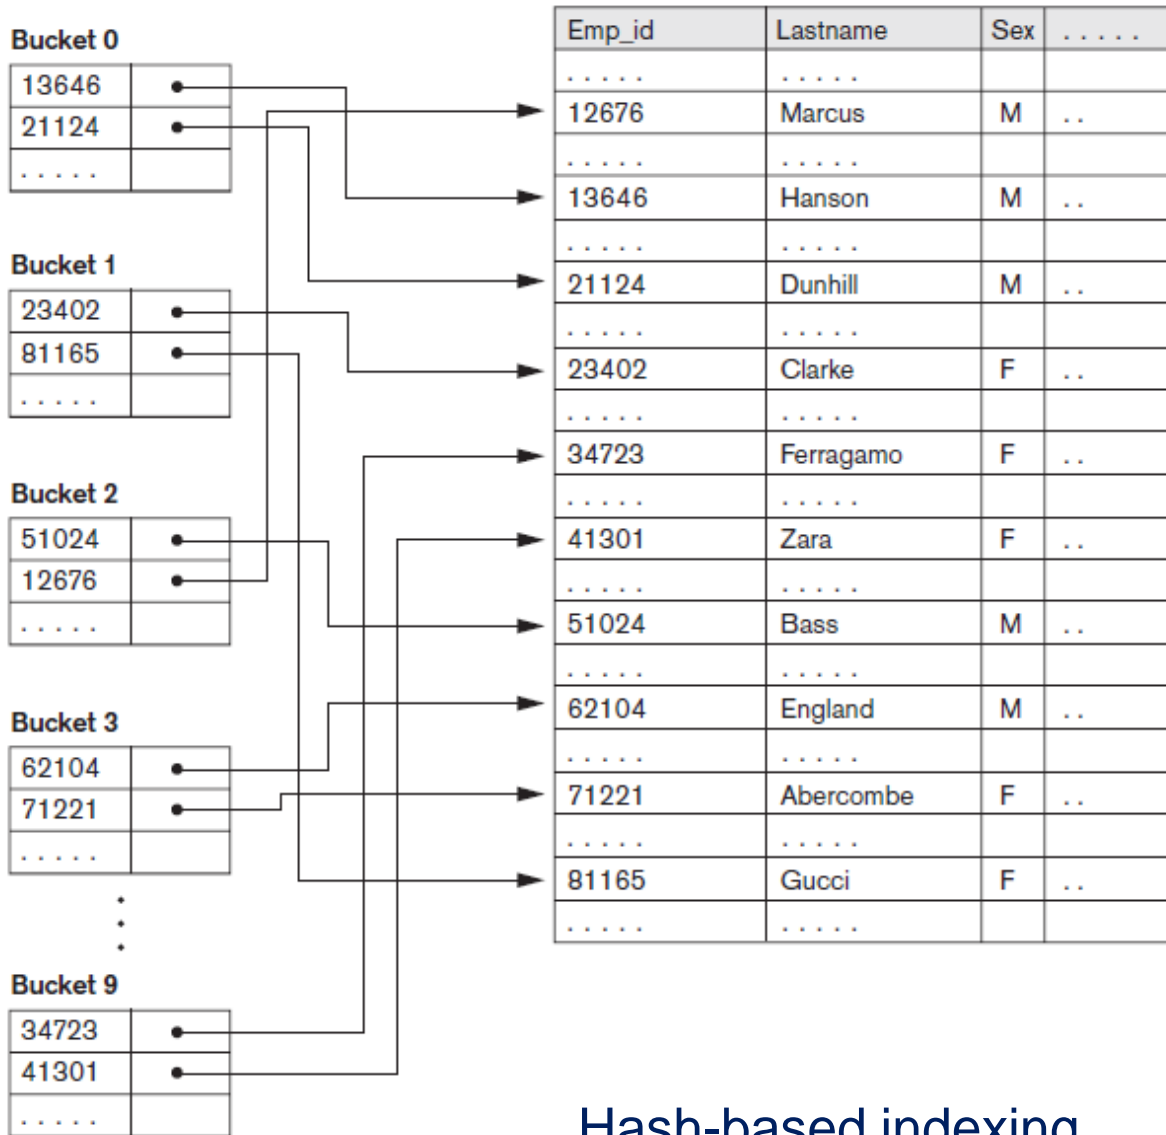


Example of a grid array on Dno and Age attributes

Other Types of Indexes

- Hash indexes
 - Secondary structure for file access
 - Uses hashing on a search key other than the one used for the primary data file organization
 - Index entries of form (K, P_r) or (K, P)
 - P_r : pointer to the record containing the key
 - P : pointer to the block (or bucket) containing the record for that key

Hash Indexes (cont'd.)



Some General Issues Concerning Indexing

- **Physical index**
 - Pointer specifies physical record address
 - **Disadvantage**: pointer must be changed if record is moved
- **Logical index**
 - Used when physical record addresses expected to change frequently
 - Entries of the form (K, K_p)

Tuning Indexes

- Tuning goals
 - Dynamically **evaluate** requirements
 - **Reorganize indexes** to yield best performance
- **Reasons for revising** initial index choice
 - Certain queries may take too long to run due to lack of an index
 - Certain indexes may not get utilized
 - Certain indexes may undergo too much updating if based on an attribute that undergoes frequent changes

Physical Database Design in Relational Databases

- Physical design goals
 - Create appropriate structure for data in storage
 - Guarantee good performance
- Must know job mix for a particular set of database system applications
- Analyzing the database queries and transactions
 - Information about each retrieval query
 - Information about each update transaction

Physical Database Design in Relational Databases (cont'd.)

- Analyzing the expected frequency of invocation of queries and transactions
 - Expected frequency of using each attribute as a selection or join attribute
 - 80-20 rule: 80 percent of processing accounted for by only 20 percent of queries and transactions
- Analyzing the time constraints of queries and transactions
 - Selection attributes associated with time constraints are candidates for primary access structures

Physical Database Design Decisions

- Design decisions about indexing
 - Whether to index an attribute
 - Attribute is a key or used by a query
 - What attribute(s) to index on
 - Single or multiple
 - Whether to set up a clustered index
 - One per table
 - Whether to use a hash index over a tree index
 - Hash indexes do not support range queries
 - Whether to use dynamic hashing
 - Appropriate for very volatile files

Summary

- Indexes are access structures that improve efficiency of record retrieval from a data file
- Ordered single-level index types
 - Primary, clustering, and secondary
- Multilevel indexes can be implemented as B-trees and B+ -trees
 - Dynamic structures
- Multiple key access methods
- Logical and physical indexes

Summary (Cont.)

- Tree-structured indexes are ideal for range-searches, also good for equality searches.
- B+ tree is a dynamic structure.
 - Inserts/deletes leave tree height-balanced; High fanout (**F**) means depth rarely more than 3 or 4.
 - Almost always much better than maintaining a sorted file.
 - Typically, 67% occupancy on average.
 - If data entries are data records, splits can change rids!
- Most widely used index in database management systems because of its versatility. One of the most optimized components of a DBMS.